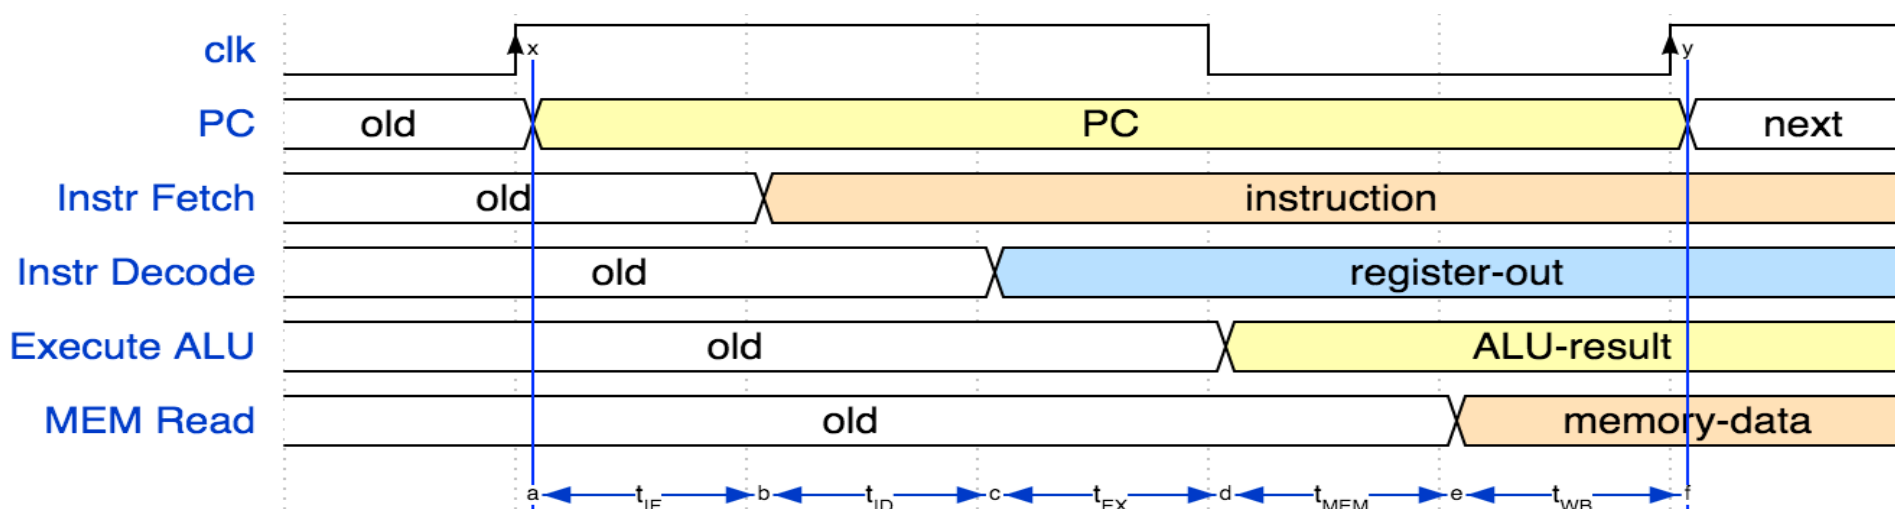


COMPUTER ORGANIZATION

Lecture 9 Pipeline Overview

2024 Spring

Instruction Timing – Single Cycle



Theoretical modeling(ideal single cycle), which could be slightly different with lab implementation

Example:

IF	ID	EX	MEM	WB	Total
I-MEM	Reg Read	ALU	D-MEM	Reg W	
200 ps	100 ps	200 ps	200 ps	100 ps	800 ps

1. IF: Instruction fetch from memory
2. ID: Instruction decode & register read
3. EX: Execute operation or calculate address
4. MEM: Access memory operand
5. WB: Write result back to register


Performance Issues

Instr	IF = 200ps	ID = 100ps	ALU = 200ps	MEM=200ps	WB = 100ps	Total
add	√	√	√		√	600ps
beq	√	√	√			500ps
lw	√	√	√	√	√	800ps
sw	√	√	√	√		700ps

- Longest delay determines clock period
 - Critical path: load instruction
 - Instruction memory → register file → ALU → data memory → register file
 - Maximum clock rate in the above example
 - $f_{\max} = 1/800\text{ps} = 1.25 \text{ GHz}$
- Most blocks idle most of the time
 - E.g.: How can we keep ALU busy all the time?
 - Idea: Factories use three employee shifts - equipment is always busy! i.e., Pipelining


Pipelining Analogy

- Pipelined laundry: overlapping execution
 - Parallelism improves performance



- Four loads(负载):
 - Speedup

$$= 8 / 3.5 = 2.3$$



- n tasks:
 - Speedup

$$= 2n / (1.5 + 0.5n) \approx 4$$

$$= \text{number of stages}$$

Pipeline Performance

Pipeline Speedup

	Single Cycle	Pipelining
Timing	$t_{stage} = 100/200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, $freq$	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Speedup	1 x	4 x

$$\text{Time between instructions}_{\text{pipelined}} = \frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$$

- If all stages are balanced
 - i.e., all take the same time
- If not balanced, speedup is less
- Speedup due to **increased throughput**
 - Latency (time for each instruction) **does not decrease**



Pipelining and ISA Design

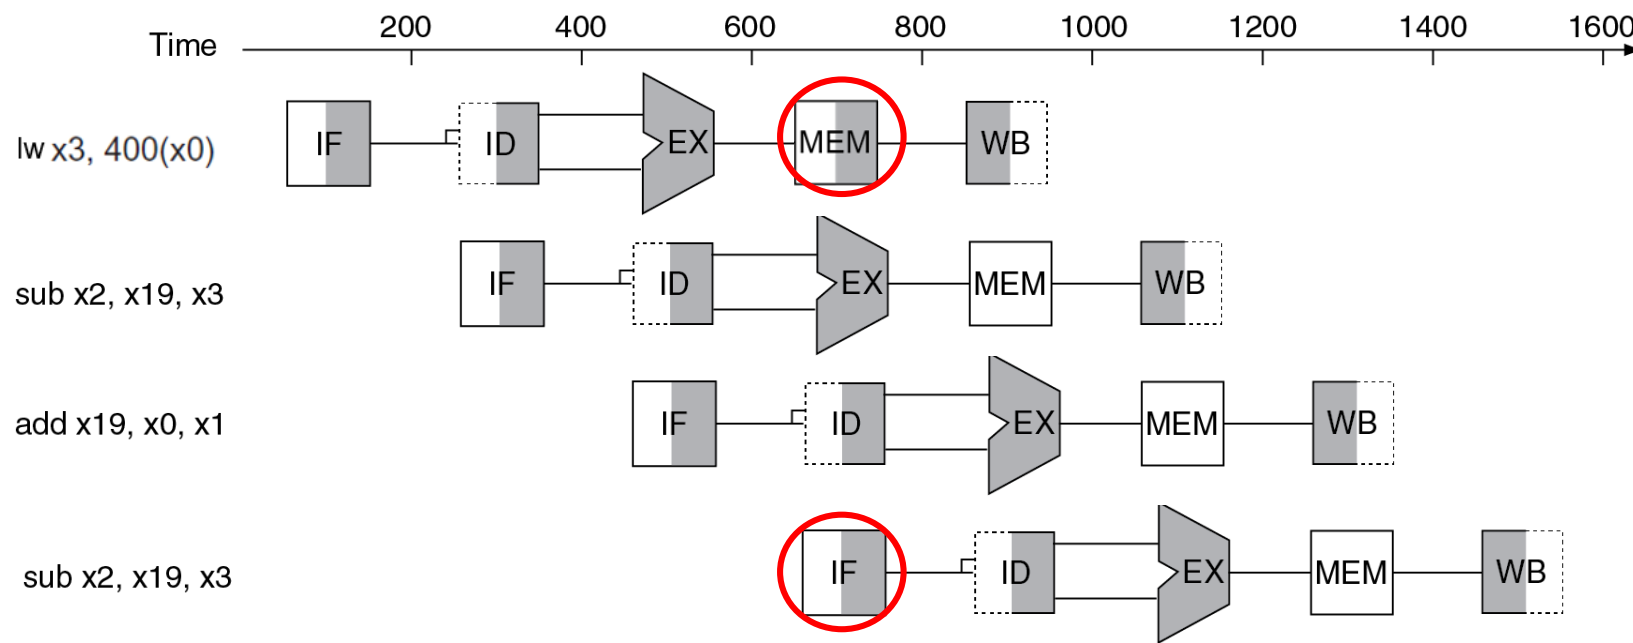
- RISC-V ISA designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage

Hazards

- Situations that prevent starting the next instruction in the next cycle
 - **Structure hazards**
 - A required resource is busy
 - **Data hazard**
 - Need to wait for previous instruction to complete its data read/write
 - **Control hazard**
 - Deciding on control action depends on previous instruction
- Can usually resolve hazards by **stall (waiting)**
 - pipeline control must detect the hazard
 - and take action to resolve hazards

Structure Hazards

- Conflict for use of a resource
- Instruction and data memory used simultaneously
- In RISC-V, use two separate memories
 - Instruction Mem and Data Mem



Data Hazards

- When data needed to execute the instruction is not yet available.
 - Register usage
 - Load-Use
- Solution
 - Stall
 - Forwarding
 - Stall + Forwarding
 - Code scheduling

Resolve Data Hazards 1: Stall

- Inserting NOP instructions (stall/bubble)

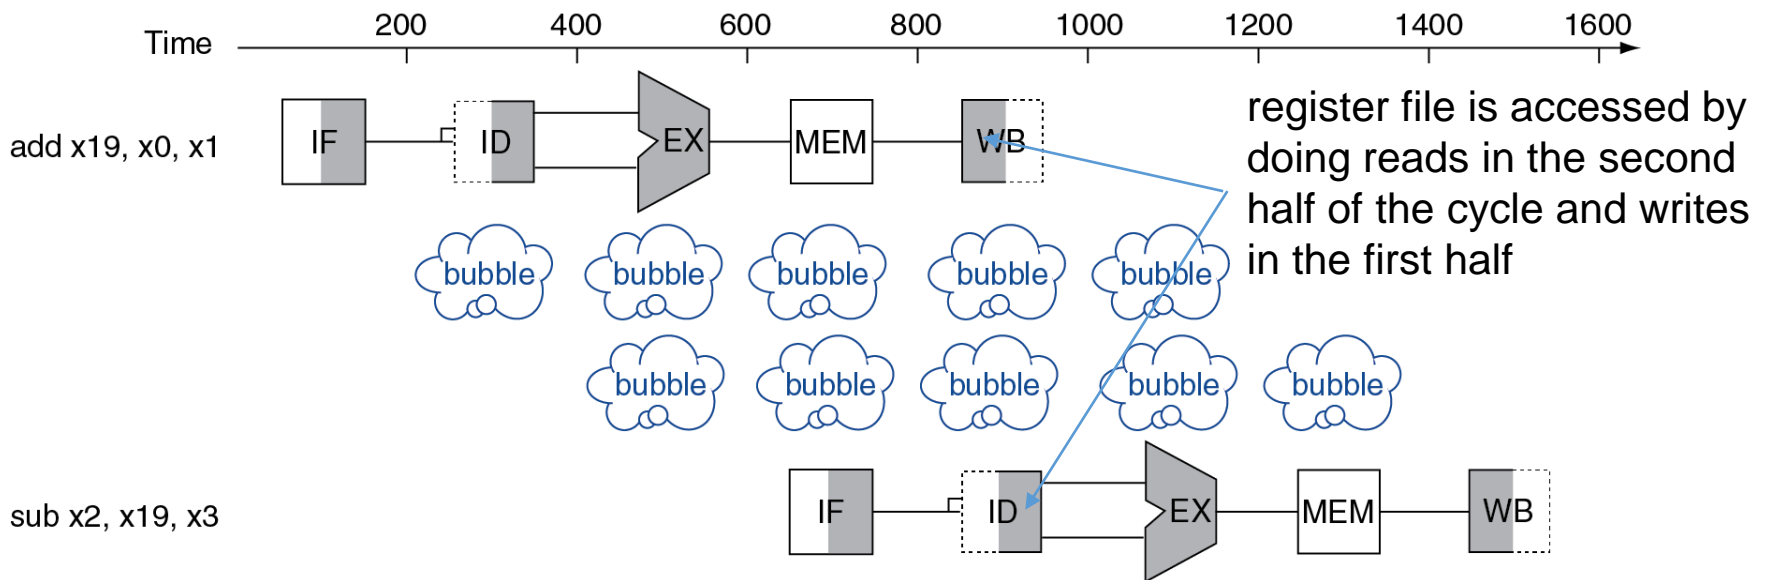
- add **x19**, x0, x1
 - sub x2, **x19**, x3



- 3 stalls is also ok

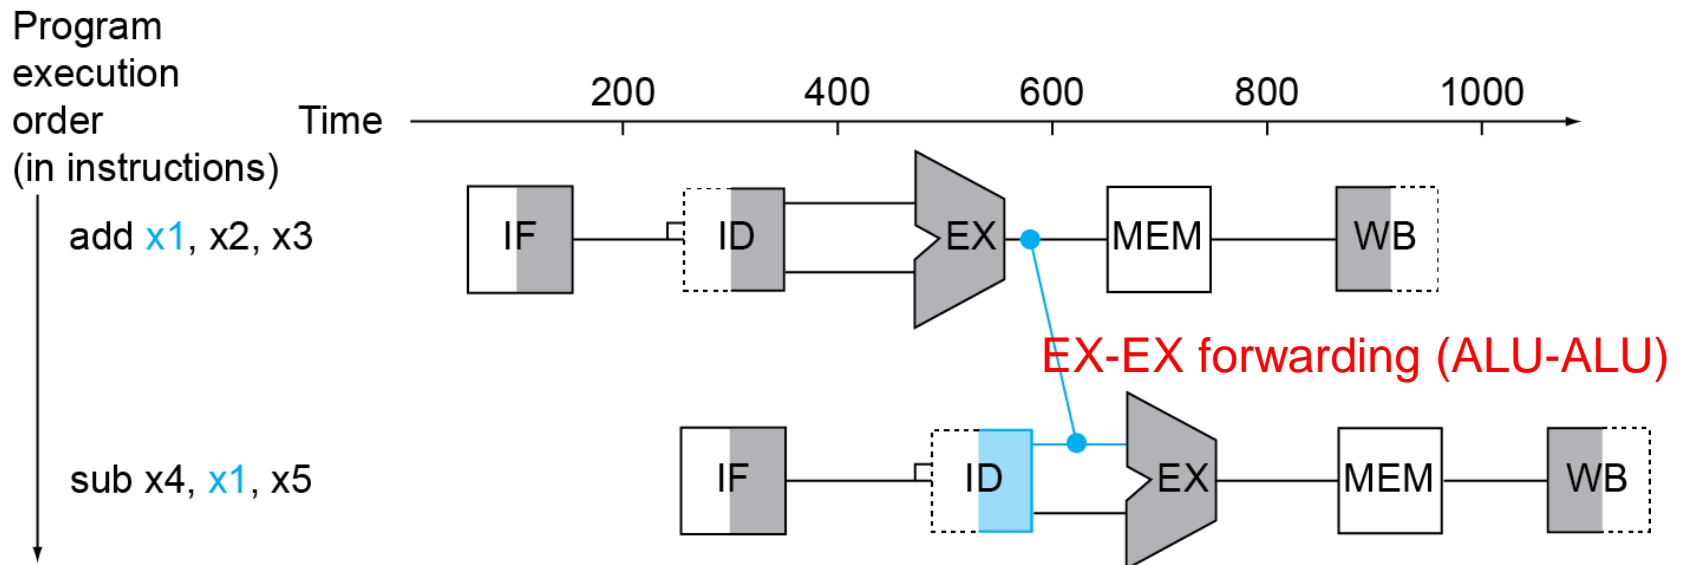


add **x19**, x0, x1
 NOP
 NOP
 sub x2, **x19**, x3



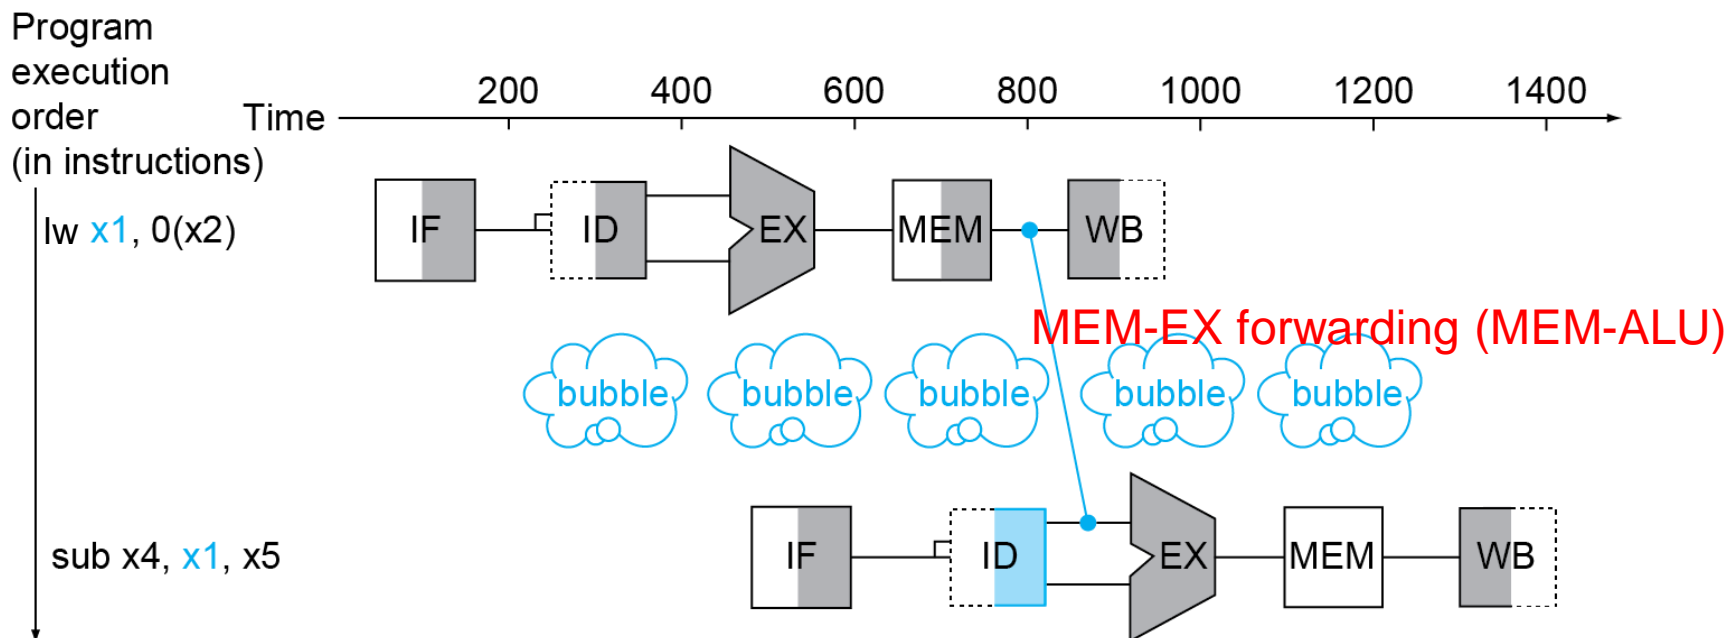
Resolve Data Hazards 2: Forwarding

- **Forwarding** can help to solve data hazard
- Core idea: Use result immediately when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath
 - Add a **bypassing line** to connect the output of EX to the input



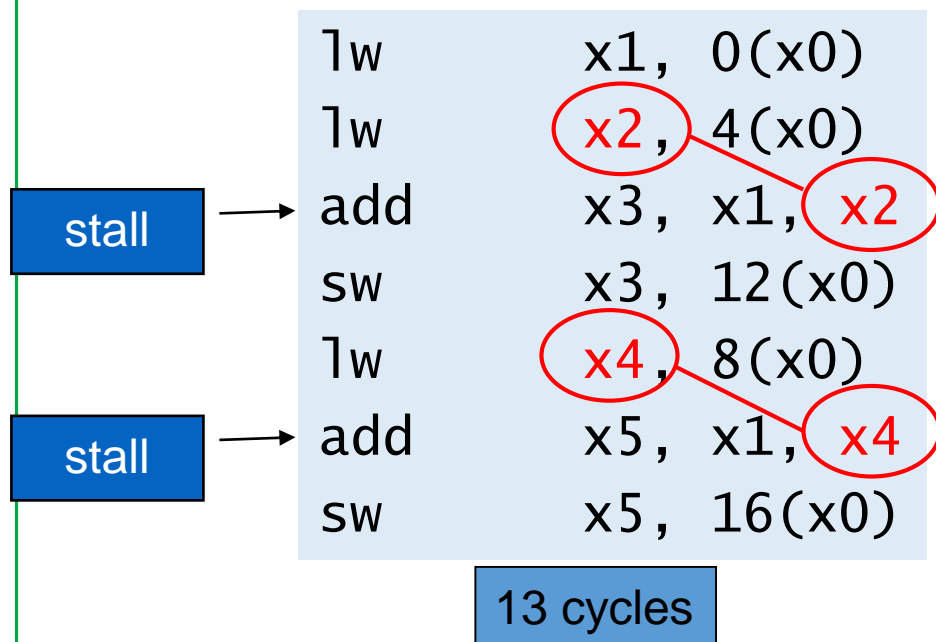
Resolve Data Hazards 3: Stall + forwarding

- For Load-Use Data Hazard, Can't avoid stalls by forwarding
 - If value not computed when needed
 - Can't forward backward in time!
 - E.g. lw **x1**, 0(x2) sub x4, **x1**, x5

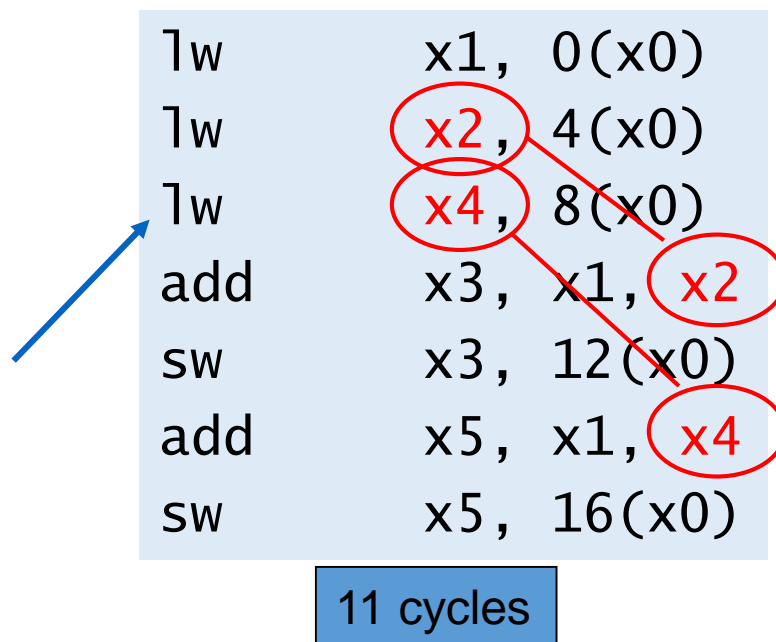


Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for $a = b + e$; $c = b + f$;
- (x1:b, x2:e, x3:a, x4:f, x5:c)



Assume forwarding is available, we still need 2 stalls to resolve data hazard



Reorder code to avoid stalls



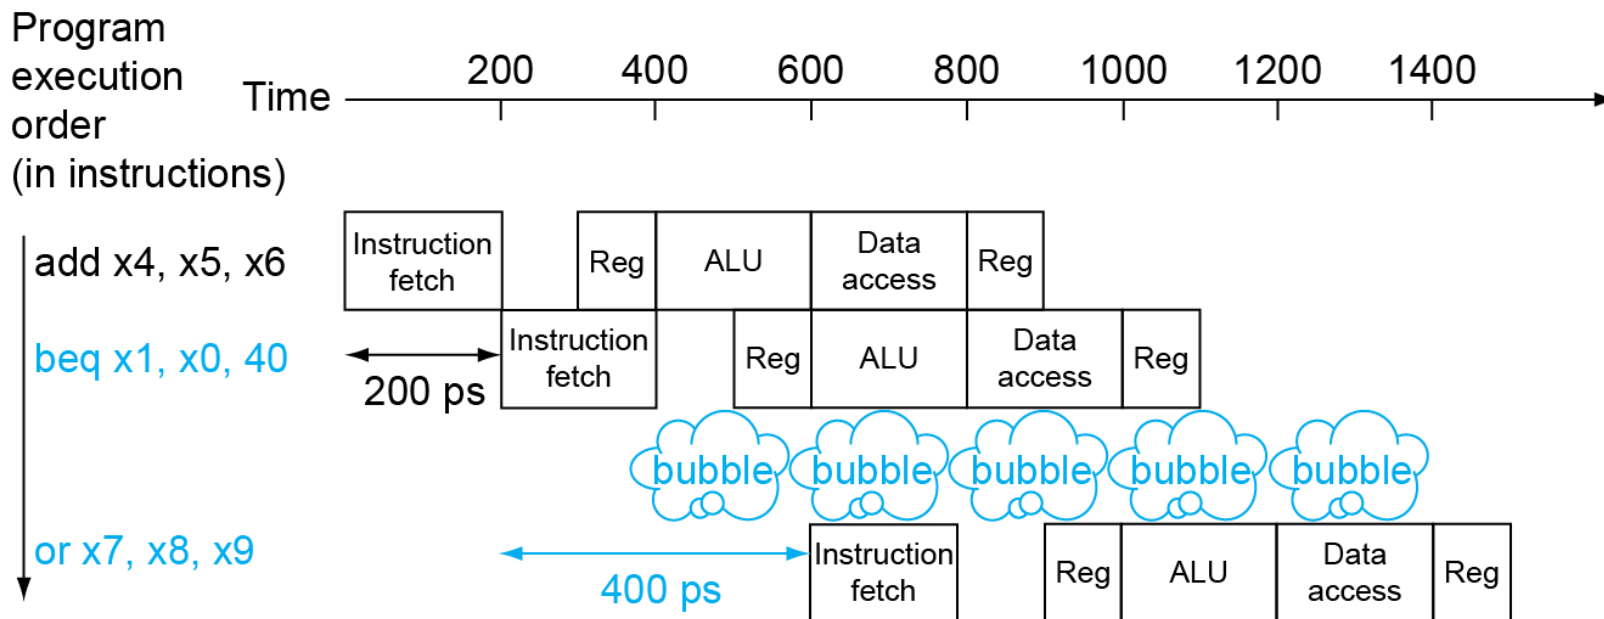
Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch
- In RISC-V pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage



Resolve Control Hazard 1: Stall on Branch

- Wait until branch outcome determined before fetching next instruction
 - Note: we assume extra hardware is added to determine the branch outcome early in ID stage





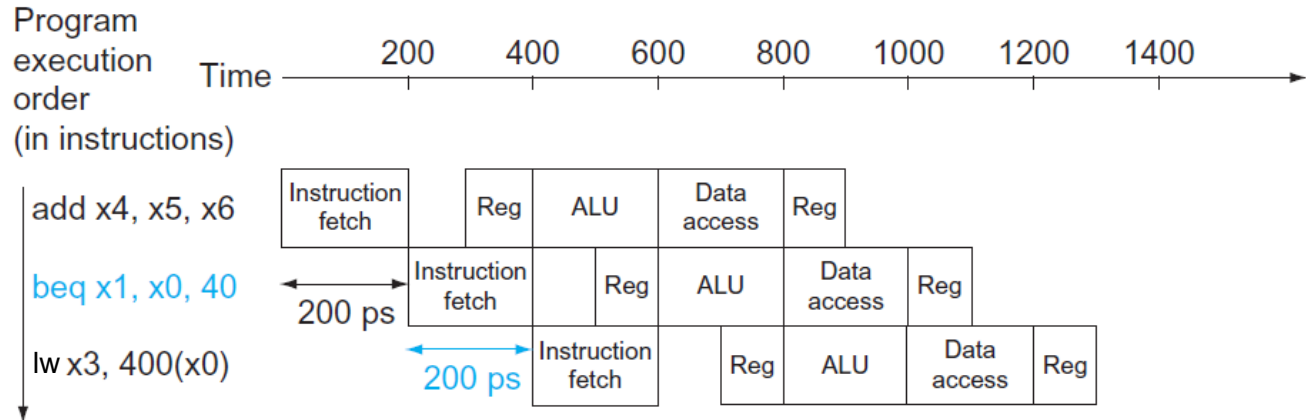
Resolve Control Hazard 2: Branch Prediction

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong
- In RISC-V pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay

RISC-V with Predict Not Taken


Prediction correct

no stall



Prediction incorrect

one stall



For Internal Use Only!

More-Realistic Branch Prediction

- Static branch prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history



Pipeline Summary

- Pipelining improves performance by increasing instruction **throughput**
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- Subject to hazards
 - Structure, data, control
- Instruction set design affects complexity of pipeline implementation