

Adversarial Search

Minimax Algorithm

- **Minimax Algorithm** is a basic method to solve game-tree problems.

MINIMAX-DECISION and EXPECTIMINIMAX

function MINIMAX-DECISION(*state*) **returns** *an action*
 return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\text{state}, a))$


function MAX-VALUE(*state*) **returns** *a utility value*
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(\text{state}, a)))$
 return *v*

function MIN-VALUE(*state*) **returns** *a utility value*
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow \infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(\text{state}, a)))$
 return *v*

function EXPECTIMINIMAX(*s*) =
 UTILITY(*s*) **if** TERMINAL-TEST(*s*)
 $\max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a))$ **if** PLAYER(*s*) = MAX
 $\min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a))$ **if** PLAYER(*s*) = MIN
 $\sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r))$ **if** PLAYER(*s*) = CHANCE


A Basic Example

- ▶ Fill X or O in the box
- ▶ X plays first
- ▶ Termination:
 - If there are two consecutive X or O, then it wins
- ▶ Utility function:
 - If X wins, score=1
 - If O wins, score=-1
 - If tie, score=0




A Basic Example: The First Level with Max

MAX(): If X wants to win, he/she will choose the strategy with the largest score





A Basic Example: The Second Level with Min

MIN(): If X has chosen the center-above position, the opponent O would choose the strategy with **the smallest score**.




A Basic Example: Third Level to Sixth Level




Fill in the Blank


Please complete the scores in the remaining boxes according to the above procedure

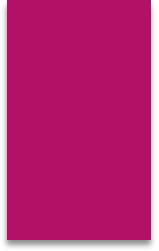


Fill in the Blank



Which position would X choose in the first step?






In the implementation of the Minimax algorithm:

- ▶ If we arrive at a symmetric case, is it necessary to search twice?
- ▶ In the process, can you summarize which searches are redundant?
- ▶ If you were to design a suitable evaluation function, how would you design it?


Pruning

- ▶ The Level-3 needs to get the maximum value of the Level-4. As the search proceed, the leftmost node in Level-4 already gets value 1 and return it to Level-3. Then, Level-3 continue to call the second node in Level-4, which already gets value -1 . Note $-1 < 1$.
- ▶ Then, does the second node in Level-4 need to continue to search its right branch?
- ▶ If the right branch gets value larger than -1 , it is obvious that Min still gets the value -1 when the search is finished. If the right branch gets value smaller than -1 , it is clear that Max in Level-3 still gets the value 1. So, the result of Level-3 will remain the same no matter which value the right branch of the second node in Level-4 returns.
- ▶ Conclusion: The second right branch of the fourth layer can be cut off.




Alpha Pruning

- ▶ The value of **node A** should be the greater of the values of **node B** and **node C**. **Node B** is now known to have a value greater than the value of **node D**. Since the value of **node C** should be the **smallest** of the values of its children, this minimum value must be no larger than the value of **node D**, and therefore must be less than the value of **node B**, indicating the meaningless of the search of other children of **node C**, e.g., **node E** and **F**. Now, we can cut off the subtree rooted at **node C**. This optimization is called **Alpha pruning**.
- ▶ Question: What happens if searching **branches E** and **F** are in front of **D**?




Alpha Pruning: Explanation

- ▶ **MAX Level:** At node A, the maximum value found in the child node is saved in **alpha**. This **alpha** value is passed to the next level along with the function call.
- ▶ The next level is **MIN level**. The minimum value currently found by the node of the **MIN level** is no larger than **alpha** value, so there is no need to continue searching.
- ▶ A sub-node needs to keep updating its own **beta** value. If the node branch does not terminate, and the currently found minimum value $<$ **beta**, we need to update the **beta** value and pass it to the next level of the node.




Beta Pruning

- ▶ The value of **node A** should be the lesser of the values of **node B** and **node C**. **Node B** is known to have a value less than the value of **node D**. Since the value of **node C** should be the **largest** of its sub-node values, this maximum value must be no less than the value of **node D**, and therefore greater than the value of **node B**, indicating that continuing to search for other children of **node C** have no meaning, and all subtrees rooted at **node C** can be **cut off**. This optimization is called **Beta pruning**.
- ▶ Question: What would happen if the branches of **E** and **F** are in front of **D**?



Beta Pruning: Explanation

- **MIN Level:** At **node A**, the **minimum value** found in the child node is saved in **beta**. This **beta** value is passed to the next level along with the function call.
- The next level is **MAX level**. The maximum value currently found by the node of the **MAX level** is no less than **beta** value, so there is no need to continue searching.
- A sub-node (max) needs to keep updating its own **alpha** value. If the node branch does not terminate, and the currently found maximum value $>$ **alpha**, we need to **update** the **alpha** value and pass it to the next level of the node.



Alpha-Beta Pruning

- ▶ Applying Alpha-Beta pruning to the Minimax algorithm, we derive the Alpha-Beta search algorithm.
- ▶ Its optimization uses properties of Minimax and does not change the result of Minimax.
- ▶ The optimization depends on the order of nodes.


```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(\text{state}, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$  Beta Pruning  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$  update alpha value  
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(\text{state}, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$  Alpha Pruning  
     $\beta \leftarrow \text{MIN}(\beta, v)$  Update beta value  
  return  $v$ 
```


Application of Alpha-Beta Pruning

- ▶ The execution result of a MINIMAX algorithm is shown in the right figure.
- ▶ Using the Alpha-Beta pruning algorithm to prune the right figure.



Tic-Tac-Toe

► <http://aimacode.github.io/aima-javascript/5-Adversarial-Search/>