

Principles of Database Systems (CS307)

Lecture 13: Transaction

Ran Cheng

Department of Computer Science and Engineering
Southern University of Science and Technology

- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7th Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.

Transaction in Real Life

- “An exchange of goods for money”
 - A series of steps
 - All or nothing



Flickr:BracketingLife (Clarence)



Transaction in Computer

- A **transaction** is a unit of program execution that accesses and possibly updates various data items
 - A classical example in database: money transfer

E.g., transaction to transfer CNY ¥50 from account A to account B:

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

Transaction in Computer

- A **transaction** is a unit of program execution that accesses and possibly updates various data items
 - A classical example in database: money transfer

E.g., transaction to transfer CNY ¥50 from account A to account B:


```
1. read(A)
2. A := A - 50
3. write(A)
4. read(B)
5. B := B + 50
6. write(B)
```

- Two main issues to deal with:
 - **Failures** of various kinds, such as hardware failures and system crashes
 - **Concurrent execution** of multiple transactions

How to perform such Transactions?

- The **BEGIN** statement starts the transaction.
- The **UPDATE** statements implicitly include the reading and writing of the account balances.
- The **COMMIT** statement finalizes the transaction, making all changes permanent.

sql

 Copy code

```
-- Start the transaction
BEGIN;

-- Step 1: Read balance from account A (This is implicit in the UPDATE command)
-- Step 2: Deduct ¥50 from account A
UPDATE accounts SET balance = balance - 50 WHERE account_id = 'A';

-- Step 3: Write the updated balance to account A (Also implicit in the UPDATE)

-- Step 4: Read balance from account B (Implicit in the UPDATE command)
-- Step 5: Add ¥50 to account B
UPDATE accounts SET balance = balance + 50 WHERE account_id = 'B';

-- Step 6: Write the updated balance to account B (Also implicit in the UPDATE)

-- Commit the transaction
COMMIT;
```

How it works? -- ACID Properties

- A **transaction** is a unit of program execution that accesses and possibly updates various data items
 - To preserve the integrity of data the database system must ensure:

Atomicity: Either all operations of the transaction are properly reflected in the database, or none are

Consistency: Execution of a transaction in isolation preserves the consistency of the database.

Isolation: Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.

- That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.

Durability: After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

Requirements in Transactions

- Atomicity Requirement
 - If the **transaction fails** after step 3 and before step 6, **money will be “lost”** leading to an **inconsistent database state**
 - Failure could be due to software or hardware
 - The system should ensure that updates of a partially executed transaction are not reflected in the database

e.g., transaction to transfer CNY ¥50 from account A to account B:

1. **read(A)**
2. **$A := A - 50$**
3. **write(A)**
4. **read(B)**
5. **$B := B + 50$**
6. **write(B)**

Requirements in Transactions

- Consistency Requirement
 - Explicitly specified integrity constraints such as primary keys and foreign keys
 - Implicit integrity constraints
 - e.g., sum of balances of all accounts
 - In the example: The sum of A and B is unchanged by the execution of the transaction

E.g., transaction to transfer CNY ¥50 from account A to account B:

```
1. read(A)
2. A := A - 50
3. write(A)
4. read(B)
5. B := B + 50
6. write(B)
```



Requirements in Transactions

- Isolaton Requirement

- If between steps 3 and 6, **another transaction T2** is allowed to access the partially updated database, it will see an inconsistent database
 - The sum $A + B$ will be less than it should be

T1	T2
1. read(A)	
2. $A := A - 50$	
3. write(A)	
	read(A), read(B), print(A+B)
4. read(B)	
5. $B := B + 50$	
6. write(B)	

- Isolation can be ensured trivially by running transactions serially, that is, one after the other
 - However, executing multiple transactions concurrently has significant benefits

Requirements in Transactions

- Durability Requirement
 - Once the user has been notified that **the transaction has completed** (i.e., the transfer of the ¥50 has taken place), **the updates to the database** by the transaction **must persist** even if there are software or hardware failures.
 - It guarantees that once a transaction has been committed, all the changes made in that transaction are **permanent** and will persist even in the event of a system failure, such as software crashes or hardware malfunctions.

Requirements in Transactions -- ACID

- A **transaction** is a unit of program execution that accesses and possibly updates various data items
 - To preserve the integrity of data the database system must ensure:

Atomicity: Either all operations of the transaction are properly reflected in the database, or none are

Consistency: Execution of a transaction in isolation preserves the consistency of the database.


Isolation: Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.

- That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.

Durability: After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.


Why ACID properties work?

- The properties are underpinned by the **Serializability Theory**, which is a key concept in transaction processing.



Transaction State

- Active
 - The initial state; the transaction stays in this state while it is executing
- Partially committed
 - After the final statement has been executed.
- Failed
 - After the discovery that normal execution can no longer proceed.
- Aborted – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
 - Restart the transaction
 - Can be done only if no internal logical error
 - Kill the transaction
- Committed
 - After successful completion.



Concurrent Executions

- **Multiple transactions** are allowed to **run concurrently** in the system.

Advantages are:

- Increased processor and disk utilization, leading to better transaction throughput
 - E.g., one transaction can be using the CPU while another is reading from or writing to the disk
- Reduced average response time for transactions
 - Short transactions do not need to wait behind long ones
- **Concurrency control schemes** – mechanisms to achieve isolation
 - That is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database

Schedules

- **Schedule** – a sequences of instructions that specify the **chronological order** (时间序列) in which instructions of concurrent transactions are executed
 - A schedule for a set of transactions must consist of **all** instructions of those transactions
 - Must preserve the **order** in which the instructions appear in each individual transaction
- A transaction that successfully completes its execution will have a **commit** instruction as the last statement
 - By default, transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement

Schedule 1

- Let T_1 transfer CNY ¥50 from A to B, and T_2 transfer 10% of the balance from A to B
- A **serial schedule** in which T_1 is followed by T_2
:

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 2

- A **serial schedule** where T_2 is followed by T_1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 3

- Let T_1 and T_2 be the transactions defined previously
 - The following schedule is not a **serial schedule**, but it is **equivalent** to Schedule 1
 - In Schedules 1, 2 and 3, the sum $A + B$ is preserved.

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Schedule 4

- The following concurrent schedule does **not** preserve the value of $(A + B)$

T_1	T_2
read (A) $A := A - 50$	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	
	$B := B + temp$ write (B) commit

Serializability

- Basic Assumption:
 - Each transaction preserves database consistency
 - Thus, serial execution of a set of transactions preserves database consistency
- A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule
 - Different forms of schedule equivalence give rise to the notions of:
 - 1. Conflict serializability
 - 2. * View serializability

Simplified View of Transactions

- We ignore operations other than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- Our simplified schedules consist of only **read** and **write** instructions.

Conflicting Instructions

- Instructions I_i and I_j , of transactions T_i and T_j respectively, **conflict** if and only if there exists any item Q accessed by both I_i and I_j , and at least one of these instructions wrote Q .
 - 1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict
 - 2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 - 3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict
 - 4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict
- Intuitively, a **conflict** between I_i and I_j forces a (logical) temporal order between them.
 - If I_i and I_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

Conflict Serializability

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are **conflict equivalent**
- We say that a schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule

Conflict Serializability

- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1
 - ... by series of swaps of **non-conflicting instructions**
 - Therefore, **Schedule 3** is *conflict serializable*.

Operations on different data

- ... and hence swappable in temporal order

T_1	T_2
read (A) write (A)	
	read (A) write (A)
	read (B) write (B)

Schedule 3

T_1	T_2
read (A) write (A) read (B) write (B)	
	read (A) write (A) read (B) write (B)

Schedule 6

Conflict Serializability

- Example of a schedule that is not conflict serializable:

T_3	T_4
read (Q)	
write (Q)	write (Q)

- We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

* View Serializability

- Let S and S' be two schedules with the same set of transactions. S and S' are **view equivalent** if the following three conditions are met, for each data item Q ,
 - If in schedule S , transaction T_i reads the initial value of Q , then in schedule S' also transaction T_i must read the initial value of Q .
 - If in schedule S transaction T_i executes **read**(Q), and that value was produced by transaction T_j (if any), then in schedule S' also transaction T_i must read the value of Q that was produced by the same **write**(Q) operation of transaction T_j .
 - The transaction (if any) that performs the final **write**(Q) operation in schedule S must also perform the final **write**(Q) operation in schedule S' .
- As can be seen, view equivalence is also based purely on **reads** and **writes** alone.

* View Serializability

- A schedule S is **view serializable** if it is view-equivalent to a serial schedule
- Every conflict serializable schedule is also view serializable
- Below is a schedule which is view-serializable but **not** conflict serializable

Two “blind writes” in T_{27} and T_{28}

- Since the written values were not used anywhere else

T_{27}	T_{28}	T_{29}
read (Q)	write (Q)	
write (Q)		
		write (Q)

Overwrites values from T_{27} and T_{28}

- ... and hence, swapping write(Q) in T_{27} and T_{28} will not affect the resulting value of Q

- What serial schedule is above equivalent to?
- Every view-serializable schedule that is not conflict serializable has **blind writes**
 - **Blind write:** Write operations without reading it before

Testing for Serializability


- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence graph**
 - A directed graph where the vertices are the transactions (names of the transactions)
 - We draw an arc from T_i to T_j if the two transactions **conflict**
 - which means, in the schedule S , T_i must appear earlier than T_j
 - We may label the arc by the item that was accessed.

Conflict – At least one of the following situations exists for a data item Q :

- $T_i: \text{write}(Q) \rightarrow T_j: \text{read}(Q)$
- $T_i: \text{read}(Q) \rightarrow T_j: \text{write}(Q)$
- $T_i: \text{write}(Q) \rightarrow T_j: \text{write}(Q)$


Testing for Serializability

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



Schedule 1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit




Schedule 2

Testing for Serializability

T_1	T_2
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	$B := B + temp$ write (B) commit

Schedule 4






Testing for Serializability

- A schedule is conflict serializable if and only if its precedence graph is **acyclic**

Cycle-detection: You will learn it in your DSAA course.

- If the precedence graph is acyclic, the serializability order can be obtained by a **topological sorting** of the graph
 - E.g., The topological order of (a) can be (b) and (c)



Recoverable Schedules

- Need to address the effect of transaction failures on concurrently running transactions
- **Recoverable schedule** — if a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i appears before the commit operation of T_j .
- The following schedule is **not** recoverable

T_8	T_9
read (A) write (A)	
	read (A) commit
read (B)	

- If T_8 should abort, T_9 would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.

Weak Levels of Consistency

- Some applications are willing to live with **weak levels of consistency**, allowing schedules that are not serializable
 - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - Such transactions do not need to be serializable with respect to other transactions
- Purpose: **Trade-off** between accuracy and performance

Levels of Consistency (in SQL-92)

- Serializable (Strongest)
 - Default
- Repeatable read — only committed records to be read.
 - Repeated reads of same record must return same value.
 - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- Read committed — only committed records can be read.
 - Successive reads of record may return different (but committed) values.
- Read uncommitted (Weakest) — even uncommitted records may be read.

Levels of Consistency

- Lower degrees of consistency can be useful for gathering approximate information about the database
- **Warning:** some database systems do not ensure serializable schedules by default
 - E.g., Oracle (and PostgreSQL prior to version 9) by default support a level of consistency called **snapshot isolation** (not part of the SQL standard)
- **Warning 2:** All SQL-92 consistency levels infer that dirty writes are prohibited
 - **Dirty write** - when one transaction overwrites a value that has previously been written by another still in-flight transaction