

Replicating Side Channel Attacks on RISC-V

Ben Chen, Shuwei Zhang, Jikun Liao

Dept of Computer Science and Technology, SUSTech

December 26, 2024

Introduction

In this project, we

- ▶ spent 5 weeks to survey the microarchitectural vulnerabilities

Introduction

In this project, we

- ▶ spent 5 weeks to survey the microarchitectural vulnerabilities
- ▶ replicated Spectre, Phantom and Downfall on x86-64 CPU

Introduction

In this project, we

- ▶ spent 5 weeks to survey the microarchitectural vulnerabilities
- ▶ replicated Spectre, Phantom and Downfall on x86-64 CPU
- ▶ analyzed the attack surfaces on XiangShan

Introduction

In this project, we

- ▶ spent 5 weeks to survey the microarchitectural vulnerabilities
- ▶ replicated Spectre, Phantom and Downfall on x86-64 CPU
- ▶ analyzed the attack surfaces on XiangShan
- ▶ set up the workflow for running XiangShan on simulator

In this project, we

- ▶ spent 5 weeks to survey the microarchitectural vulnerabilities
- ▶ replicated Spectre, Phantom and Downfall on x86-64 CPU
- ▶ analyzed the attack surfaces on XiangShan
- ▶ set up the workflow for running XiangShan on simulator
- ▶ migrated the attacks on XiangShan

Background

Side-Channel Attacks

Side channel attacks exploit unintended information leakage in

- ▶ energy: voltage, power, radiation...

Background

Side-Channel Attacks

Side channel attacks exploit unintended information leakage in

- ▶ energy: voltage, power, radiation...
- ▶ timing: execution cost \rightarrow branching result

Background

Side-Channel Attacks

Side channel attacks exploit unintended information leakage in

- ▶ energy: voltage, power, radiation...
- ▶ timing: execution cost \rightarrow branching result
- ▶ cache: timing difference \rightarrow data present or not in cache

Background

Side-Channel Attacks

Side channel attacks exploit unintended information leakage in

- ▶ energy: voltage, power, radiation...
- ▶ timing: execution cost → branching result
- ▶ cache: timing difference → data present or not in cache
- ▶ contention: cross-core interrupts, bus, and other resource contention → network fingerprints

Background

Side-Channel Attacks

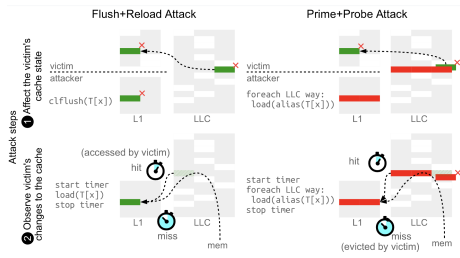
Side channel attacks exploit unintended information leakage in

- ▶ energy: voltage, power, radiation...
- ▶ timing: execution cost → branching result
- ▶ cache: timing difference → data present or not in cache
- ▶ contention: cross-core interrupts, bus, and other resource contention → network fingerprints
- ▶ other architecture: MWAIT[1] status change encoded in user-readable register

Background

Cache Side-Channel

- ▶ **Flush+Reload:** Flush all cache of attacker's probe array, encode victim's secret in offset of probe array and reload to time the difference
- ▶ **Prime+Probe:** Fill the cache with attacker's probe array, and victim's cache evict a element of probe array in cache, so attacker can tell which one is evicted



We've seen plentiful defense on RISC-V

- ▶ SafeSpec[2]: Blocking unsafe loads from altering the data cache
- ▶ SpectreGuard[3]/SpecTerminator[4]: Marking the unsafe load to prevent speculative load

but few attack on RISC-V and especially on XiangShan

- ▶ A Secure RISC[5]: Attack I\$ on C906 & U74

XiangShan is a open-source RISC-V core IP

- ▶ developed by ICT of CAS, led by Yungang Bao

XiangShan is a open-source RISC-V core IP

- ▶ developed by ICT of CAS, led by Yungang Bao
- ▶ out of order, superscalar

XiangShan is a open-source RISC-V core IP

- ▶ developed by ICT of CAS, led by Yungang Bao
- ▶ out of order, superscalar
- ▶ runs RISC-V RV64GCV, Linux-capable, FPGA-capable

XiangShan is a open-source RISC-V core IP

- ▶ developed by ICT of CAS, led by Yungang Bao
- ▶ out of order, superscalar
- ▶ runs RISC-V RV64GCV, Linux-capable, FPGA-capable
- ▶ written in Chisel, agile development and verification

XiangShan is a open-source RISC-V core IP

- ▶ developed by ICT of CAS, led by Yungang Bao
- ▶ out of order, superscalar
- ▶ runs RISC-V RV64GCV, Linux-capable, FPGA-capable
- ▶ written in Chisel, agile development and verification

Is XiangShan vulnerable to Spectre or so?

Vulnerable code in victim

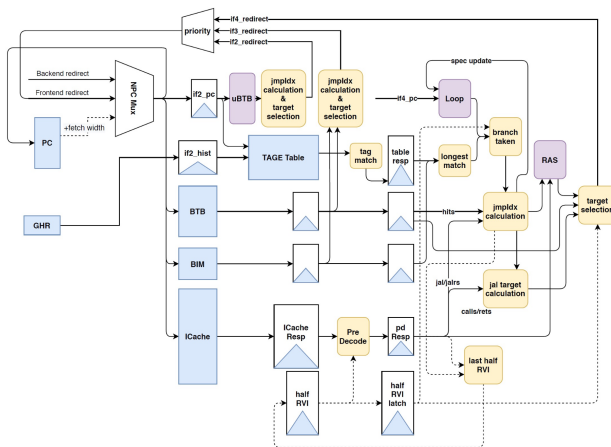
```
uint8_t array[160] = {1, 2, ..., 16};  
char* secret = "Secret goes here!";  
  
void victim_function(size_t x) {  
    if (x < array1_size) { // array1_size = 16  
        temp &= array2[array1[x] * CACHELINE_SIZE];  
    }  
}
```

Training the branch predictor

```
for (int x = 0; x < ENTRY_SIZE; ++x) {
    victim_function(x);
}
flush_cache(array2);
victim_function(secret[i++]);

for (int i = 0; i < 256; i++) {
    addr = &array2[i * CACHELINE_SIZE];
    time1 = __rdtscp(&junk); /* READ TIMER */
    junk = * addr; /* MEMORY ACCESS TO TIME */
    time2 = __rdtscp(&junk) - time1;
    /* READ TIMER & COMPUTE ELAPSED TIME */
    if (time2 <= CACHE_HIT_THRESHOLD)
        results[i]++;
}
```

XiangShan Branch Predictor



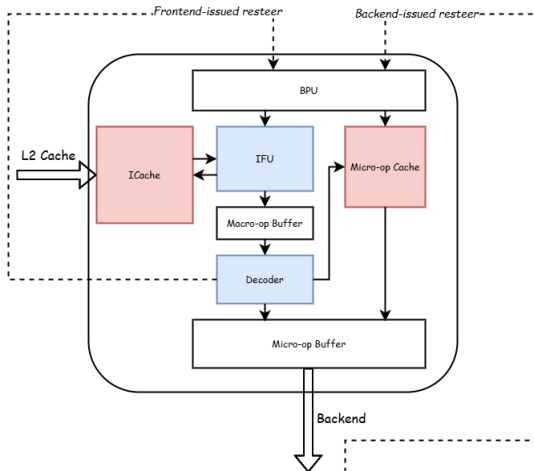
- ▶ Branch predictors: μ BTB, BTB, TAGE, RAS and loop predictor
TAGE + loop predictor \rightarrow disable the original attack
- ▶ Speculative Execution
Out-of-order, multiple issue \rightarrow Spectre-vulnerable
- ▶ Cache: L1 \$ L2
8 ways, 256 sets, 64B each line \rightarrow enough to encode a character
- ▶ Timer and cache manipulation
`rdcycle` cycle-level timer, `fence.i` memory barrier

Spectre on Nanhu

```
for (int i = 0; i <= ENTRY_SIZE; ++i) {
    int x = i < ENTRY_SIZE ? i : secret[i++]
    \\ optimize branch to jump table
    flush_cache(array2);
    victim_function(x);
}

for (int i = 0; i < 256; i++) {
    addr = &array2[i * CACHELINE_SIZE];
    time1 = rdcycle(); /* READ TIMER */
    junk = * addr; /* MEMORY ACCESS TO TIME */
    time2 = rdcycle() - time1;
    /* READ TIMER & COMPUTE ELAPSED TIME */
    if (time2 <= CACHE_HIT_THRESHOLD)
        results[i]++;
}
```

Decode Pipeline



Branch Prediction

- ▶ Prediction on instruction type.
- ▶ Prediction on branch target.

Misprediction Rester

- ▶ Frontend Rester:
 - ▶ Mismatch of predicted instruction types.
 - ▶ Incorrect branch prediction address (Direct branch).
- ▶ Backend Rester:
 - ▶ Taken/Not-taken conditional branch.
 - ▶ Incorrect branch prediction address (Indirect branch).

Phantom


Phantom on x86

Phantom

- ▶ Exploit transient window caused by frontend resteer.

Phantom Workflow

- ▶ Train A with direct / indirect branch to C.
- ▶ Execute B at aliased address to trigger misprediction to C.
- ▶ Set up observation channel to monitor C's advance in pipeline.




Phantom

Phantom on x86

IF Channel



ID Channel



EX Channel



Example: IF Channel

```
for (int j = 0; j < 8; j++)  
    train_func();  
memory_barrier //asm volatile("sfence;\nmfence;\nlfence")  
clflush((void*)monitor_addr);  
memory_barrier  
victim_func();  
delayloop(100000);  
memory_barrier  
uint64_t tim = memaccesstime((void*)monitor_addr);
```

Hardware Performance Counter used in ID Channel

- ▶ Zen2

- ▶ *de_dis_uops_from_decoder.de_dis_uops_from_opcache*

- ▶ *de_dis_uops_from_decoder.de_dis_uops_from_both*

- ▶ Zen3


- ▶ *op_cache_hit_miss.op_cache_miss*

Phantom

Phantom on x86

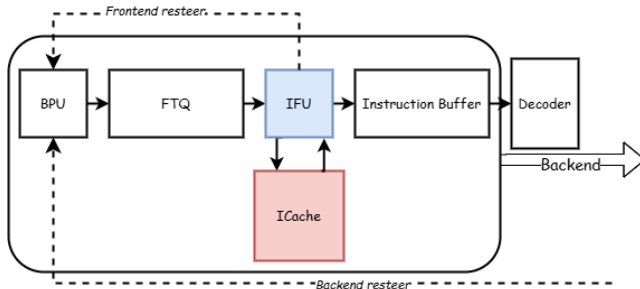
Trigger misprediction

- ▶ Create aliased address:
Flip 19th bit and 31st bit on Zen2, 21st bit and 33rd bit on Zen3.
- ▶ Set up branch history:
4~8 direct jumps separated by 128 bytes.



$$\begin{aligned} f_0 &= b_{47} \oplus b_{35} \oplus b_{23} & f_1 &= b_{47} \oplus b_{36} \oplus b_{24} \oplus b_{12} \\ f_2 &= b_{47} \oplus b_{37} \oplus b_{25} \oplus b_{13} & f_3 &= b_{47} \oplus b_{38} \oplus b_{26} \oplus b_{14} \\ f_4 &= b_{47} \oplus b_{39} \oplus b_{26} \oplus b_{13} & f_5 &= b_{47} \oplus b_{39} \oplus b_{27} \oplus b_{15} \\ f_6 &= b_{47} \oplus b_{40} \oplus b_{28} \oplus b_{16} & f_7 &= b_{47} \oplus b_{41} \oplus b_{29} \oplus b_{17} \\ f_8 &= b_{47} \oplus b_{42} \oplus b_{30} \oplus b_{18} & f_9 &= b_{47} \oplus b_{43} \oplus b_{31} \oplus b_{19} \\ f_{10} &= b_{47} \oplus b_{44} \oplus b_{32} \oplus b_{20} & f_{11} &= b_{47} \oplus b_{45} \oplus b_{33} \oplus b_{21} \end{aligned}$$

Nanhu Frontend



Main Different:

- ▶ No μ op cache.
- ▶ Integrated Instruction Fetch and Prediction Check.

Observation Channel

- ▶ IFU *frontend_icache_miss_cnt*, *frontend_flush*
- ▶ Decoder *ctrlblock_decoder_utilization*

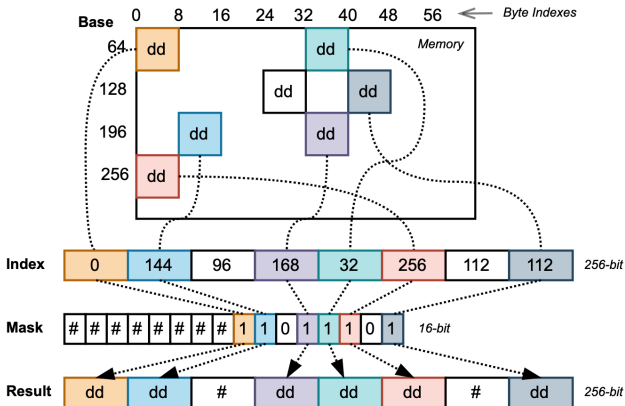
Trigger misprediction

- ▶ Create aliased address:
Using same lower 29 bits
- ▶ Set up branch history:
32 direct jumps separated by 64 bytes.

Downfall

Gather Data Sampling

Execution of gather instruction[6]



Downfall

Gather Data Sampling

Microarchitectural data sampling exploits

- ▶ buffers inside microarchitectural components like load buffer, store-commit buffer

Microarchitectural data sampling exploits

- ▶ buffers inside microarchitectural components like load buffer, store-commit buffer
- ▶ x86 has hyperthreading tech, allowing two threads run on the same core sharing the same resource → potential data stealing

Microarchitectural data sampling exploits

- ▶ buffers inside microarchitectural components like load buffer, store-commit buffer
- ▶ x86 has hyperthreading tech, allowing two threads run on the same core sharing the same resource → potential data stealing
- ▶ hard to conduct, since data in buffer vanished quickly

Microarchitectural data sampling exploits

- ▶ buffers inside microarchitectural components like load buffer, store-commit buffer
- ▶ x86 has hyperthreading tech, allowing two threads run on the same core sharing the same resource → potential data stealing
- ▶ hard to conduct, since data in buffer vanished quickly
- ▶ gather magnifies the attack by filling up the buffer with vector load

Microarchitectural data sampling exploits

- ▶ buffers inside microarchitectural components like load buffer, store-commit buffer
- ▶ x86 has hyperthreading tech, allowing two threads run on the same core sharing the same resource → potential data stealing
- ▶ hard to conduct, since data in buffer vanished quickly
- ▶ gather magnifies the attack by filling up the buffer with vector load
- ▶ then encoding

Kunminghu Architecture

Downfall

RISC-V Vector Extension

Similar to Intel's AVX and ARM's SVE

- ▶ XiangShan adds on 32 128-bit vector registers and 7 vector CSR

Downfall

RISC-V Vector Extension

Similar to Intel's AVX and ARM's SVE

- ▶ XiangShan adds on 32 128-bit vector registers and 7 vector CSR
- ▶ setting vtype and vlen before computing vectors

RISC-V Vector Extension

RISC-V Vector Extension

Similar to Intel's AVX and ARM's SVE

- ▶ XiangShan adds on 32 128-bit vector registers and 7 vector CSR
- ▶ setting vtype and vlen before computing vectors
- ▶ Vector load, indexed (gather): vlxb, vlxbu, vlxbh, vlxbhu, vlxbw, vlxbwu, vlxbd, vlxbdh, vlxbdw, vlxbd

RISC-V Vector Extension

Similar to Intel's AVX and ARM's SVE

- ▶ XiangShan adds on 32 128-bit vector registers and 7 vector CSR
- ▶ setting vtype and vlen before computing vectors
- ▶ Vector load, indexed (gather): vlxb, vlxbu, vlxbh, vlxbhu, vlxbw, vlxbwu, vlxd, vflxbh, vflxbw, vflxd
- ▶ XiangShan reuse the Load-Store Unit in execution of vector instruction → GDS attack

Downfall

Exploiting RVV Instruction

Downfall with RVV

```
fence.i
// increase the transient window
vsetvli t0, %[v1], e64, m1
// Set vector length and element width to 64 bits
vmv.v.x v0, %[mask]
// Move mask to vector register v0
vle32.v v2, (%[indices])
// Load indices into vector register v2
vluxe164.v v1, (%[src]), v2, v0.t
// Load 64-bit elements using indices and mask
vse64.v v1, (%[dst]) Store loaded elements to dst

encode_secret
flush_and_reload
```

Experiment

CPU	Generation	Memory
Xeon(R) Silver 4210	Cascade Lake	DDR4 128GB
AMD R9-3900X	Zen2	DDR3 32GB
XiangShan	Nanhu (FPGA)	DDR3 16GB
XiangShan	Kunminghu (Verilator & GEM5)	DDR3 8GB

Table: Tested machines

Experiment

CPU	Generation	Memory
Xeon(R) Silver 4210	Cascade Lake	DDR4 128GB
AMD R9-3900X	Zen2	DDR3 32GB
XiangShan	Nanhu (FPGA)	DDR3 16GB
XiangShan	Kunminghu (Verilator & GEM5)	DDR3 8GB

Table: Tested machines

CPU	Spectre	Phantom	Downfall
Xeon(R) Silver 4210	✓		✓
AMD R9-3900X		✓	
XiangShan (Nanhu)	✓	Ongoing	No RVV
XiangShan (Kunminghu)			Ongoing

Table: Result

We plan to work on..

- ▶ continue to implement the remaining attacks on XiangShan
- ▶ propose defense techniques and implement on XiangShan
- ▶ explore more attack surfaces on XiangShan and RISC-V

Conclusion

- ▶ Validated side-channel attacks on x86 Spectre v1, Phantom, and Downfall
- ▶ Theoretically proved that side-channel attacks on XiangShan and partly implement them trying to overcome the inconvenience in the ecosystem
- ▶ Continue implementing attacks on RISC-V
- ▶ Plan to mitigate the side-channels on RISC-V

- [1] Ruiyi Zhang et al. “(M)WAIT for It: Bridging the Gap between Microarchitectural and Architectural Side Channels”. In: *32nd USENIX Security Symposium (USENIX Security 23)*. Anaheim, CA: USENIX Association, Aug. 2023, pp. 7267–7284. ISBN: 978-1-939133-37-3. URL: <https://www.usenix.org/conference/usenixsecurity23/presentation/zhang-ruiyi>.
- [2] Khaled N. Khasawneh et al. “SafeSpec: Banishing the Spectre of a Meltdown with Leakage-Free Speculation”. In: *2019 56th ACM/IEEE Design Automation Conference (DAC)*. 2019, pp. 1–6.
- [3] Jacob Fustos, Farzad Farshchi, and Heechul Yun. “SpectreGuard: An Efficient Data-centric Defense Mechanism against Spectre Attacks”. In: *2019 56th ACM/IEEE Design Automation Conference (DAC)*. 2019, pp. 1–6.

- [4] Hai Jin, Zhuo He, and Weizhong Qiang. “SpecTerminator: Blocking Speculative Side Channels Based on Instruction Classes on RISC-V”. In: 20.1 (Feb. 2023). ISSN: 1544-3566. DOI: 10.1145/3566053. URL: <https://doi.org/10.1145/3566053>.
- [5] Lukas Gerlach et al. “A Security RISC: Microarchitectural Attacks on Hardware RISC-V CPUs”. In: *2023 IEEE Symposium on Security and Privacy (SP)*. 2023, pp. 2321–2338. DOI: 10.1109/SP46215.2023.10179399.
- [6] Daniel Moghimi. “Downfall: Exploiting Speculative Data Gathering”. In: *32th USENIX Security Symposium (USENIX Security 2023)*. 2023.