



CSE5014 CRYPTOGRAPHY AND NETWORK SECURITY

Dr. QI WANG

Department of Computer Science and Engineering

Office: Room413, CoE South Tower

Email: wangqi@sustech.edu.cn

Perfect security

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{x_0, x_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x_b after seeing the ciphertext $c = Enc_k(x_b)$ is at most $1/2$.



Perfect security

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{x_0, x_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses x_b after seeing the ciphertext $c = Enc_k(x_b)$ is at most $1/2$.

Theorem 1.10 (*Limitations of perfect secrecy*) There is no *perfectly secure* encryption schemes (Gen, Enc, Dec) with n -bit plaintexts and $(n - 1)$ -bit keys.

- The key is as long as the message
- Only secure if each key is used to encrypt a single message
- Trivially broken by a known-plaintext attack

ϵ -Statistical Security

- **Definition 2.1** Let X and Y be two distributions over $\{0, 1\}^n$. The *statistical distance* of X and Y , denoted by $\Delta(X, Y)$ is defined to be
$$\max_{T \subseteq \{0, 1\}^n} |\Pr[X \in T] - \Pr[Y \in T]|.$$
If $\Delta(X, Y) \leq \epsilon$, we say that $X \equiv_{\epsilon} Y$.



ϵ -Statistical Security

- **Definition 2.1** Let X and Y be two distributions over $\{0, 1\}^n$. The *statistical distance* of X and Y , denoted by $\Delta(X, Y)$ is defined to be

$$\max_{T \subseteq \{0, 1\}^n} |\Pr[X \in T] - \Pr[Y \in T]|.$$

If $\Delta(X, Y) \leq \epsilon$, we say that $X \equiv_{\epsilon} Y$.

Definition 2.2 *ϵ -Statistical Security*. An encryption scheme (Gen, Enc, Dec) is *ϵ -statistically secure* if for every pair of plaintexts m, m' , we have $Enc_{U_n}(m) \equiv_{\epsilon} Enc_{U_n}(m')$.



■ Lemma 2.3

$$\Delta(X, Y) = \frac{1}{2} \sum_{w \in \text{Supp}(X) \cup \text{Supp}(Y)} |Pr[X = w] - Pr[Y = w]|$$



■ Lemma 2.3

$$\Delta(X, Y) = \frac{1}{2} \sum_{w \in \text{Supp}(X) \cup \text{Supp}(Y)} |Pr[X = w] - Pr[Y = w]|$$

Observations:

$$0 \leq \Delta(X, Y) \leq 1$$

$$\Delta(X, Y) = 0 \text{ if } X = Y$$

$$0 \leq \Delta(X, Y) \leq \Delta(X, Z) + \Delta(Z, Y)$$

■ Lemma 2.3

$$\Delta(X, Y) = \frac{1}{2} \sum_{w \in \text{Supp}(X) \cup \text{Supp}(Y)} |Pr[X = w] - Pr[Y = w]|$$

Observations:

$$0 \leq \Delta(X, Y) \leq 1$$

$$\Delta(X, Y) = 0 \text{ if } X = Y$$

$$0 \leq \Delta(X, Y) \leq \Delta(X, Z) + \Delta(Z, Y)$$

Δ is a *metric*.

ϵ -Statistical Security

- **Lemma 2.4** Eve has at most $1/2 + \epsilon$ success probability if and only if for every pair of m_1, m_2 ,
$$\Delta(\text{Enc}_{U_n}(m_1), \text{Enc}_{U_n}(m_2)) \leq 2\epsilon.$$



ϵ -Statistical Security

- **Lemma 2.4** Eve has at most $1/2 + \epsilon$ success probability if and only if for every pair of m_1, m_2 ,
$$\Delta(\text{Enc}_{U_n}(m_1), \text{Enc}_{U_n}(m_2)) \leq 2\epsilon.$$

Theorem 2.5 Let $(\text{Gen}, \text{Enc}, \text{Dec})$ be a valid encryption with $\text{Enc} : \{0, 1\}^n \times \{0, 1\}^{n+1} \rightarrow \{0, 1\}^*$. Then there exist plaintexts m_1, m_2 with $\Delta(\text{Enc}_{U_n}(m_1), \text{Enc}_{U_n}(m_2)) > 1/2$.

Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.



Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.
 - In real life, people are using encryption with keys **shorter** than the message size to encrypt all kinds of sensitive information.
 - If the algorithm you use to break the encryption scheme runs in time 2^n , it seems **OK** since the message may be expired by then ...



Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.
 - In real life, people are using encryption with keys **shorter** than the message size to encrypt all kinds of sensitive information.
 - If the algorithm you use to break the encryption scheme runs in time 2^n , it seems **OK** since the message may be expired by then ...
- **Idea:** Would be OK if a scheme leaked information with *tiny probability* to eavesdroppers with *bounded computational resources*



Computational Security

- Statistical security does **not** allow us to break the **impossibility** result.
 - In real life, people are using encryption with keys **shorter** than the message size to encrypt all kinds of sensitive information.
 - If the algorithm you use to break the encryption scheme runs in time 2^n , it seems **OK** since the message may be expired by then ...
- **Idea:** Would be OK if a scheme leaked information with *tiny probability* to eavesdroppers with *bounded computational resources*
 - Allowing security to “**fail**” with tiny probability
 - Restricting attention to “**efficient**” attackers



Tiny probability of failure?

- Say security fails with probability 2^{-60}

Tiny probability of failure?

- Say security fails with probability 2^{-60}
 - Should we be concerned about this?



Tiny probability of failure?

- Say security fails with probability 2^{-60}
 - Should we be concerned about this?
 - With probability $> 2^{-60}$, the sender and receiver will both be struck by lightning in the next year ...
 - Something that occurs with probability $2^{-60}/\text{sec}$ is expected to occur once every **100 billion** years

Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle



Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle
 - Desktop computer $\approx 2^{57}$ keys/year
 - Supercomputer $\approx 2^{80}$ keys/year



Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle
 - Desktop computer $\approx 2^{57}$ keys/year
 - Supercomputer $\approx 2^{80}$ keys/year
 - Supercomputer since Big Bang $\approx 2^{112}$ keys
 - Restricting attention to attackers who can try 2^{112} keys is fine!

Bounded attackers?

- Consider *brute-force search* of key space; assume one key can be tested per clock cycle
 - Desktop computer $\approx 2^{57}$ keys/year
 - Supercomputer $\approx 2^{80}$ keys/year
 - Supercomputer since Big Bang $\approx 2^{112}$ keys
 - Restricting attention to attackers who can try 2^{112} keys is fine!

Modern key space: 2^{128} keys or more ...



Computational Security

- Two problems:



Computational Security

- Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).



Computational Security

- Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

e.g., the **substitution cipher** has a huge key space, **but** can be broken efficiently.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).



Computational Security

■ Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

e.g., the **substitution cipher** has a huge key space, **but** can be broken efficiently.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).

“Breaking E is very hard”?

“Problem P cannot be solved in reasonable time”?



Computational Security

■ Two problems:

1) While the particular algorithm runs in exponential time, we **cannot** guarantee that there is no other algorithm is efficient.

e.g., the **substitution cipher** has a huge key space, **but** can be broken efficiently.

2) We need a **precise** mathematical definition (like *perfect secrecy* definition).

“Breaking E is very hard”?

“Problem P cannot be solved in reasonable time”?

Q: How do we **model** the resources of Eve (the adversary)?



Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.

Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.
- Let $\Pi = (Gen, Enc, Dec)$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.
- Let $\Pi = (Gen, Enc, Dec)$ be an encryption scheme with message space \mathcal{M} , and A an adversary
Define a randomized experiment $PrivK_{A, \Pi}$:
 1. A outputs $m_0, m_1 \in \mathcal{M}$
 2. $k \leftarrow Gen, b \leftarrow \{0, 1\}, c \leftarrow Enc_k(m_b)$
 3. $b' \leftarrow A(c)$

Perfect indistinguishability

- **Definition 1.6** *Perfect secrecy*. An encryption scheme (Gen, Enc, Dec) with message space \mathcal{M} and ciphertext space \mathcal{C} is *perfectly secure* if and only if for every two distinct plaintexts $\{m_0, m_1\} \in \mathcal{M}$, and for every strategy used by Eve, if we choose at random $b \in \{0, 1\}$ and a random key $k \in \{0, 1\}^n$, then the probability that Eve guesses m_b after seeing the ciphertext $c = Enc_k(m_b)$ is at most $1/2$.
- Let $\Pi = (Gen, Enc, Dec)$ be an encryption scheme with message space \mathcal{M} , and A an adversary
Define a randomized experiment $PrivK_{A, \Pi}$:
 1. A outputs $m_0, m_1 \in \mathcal{M}$
 2. $k \leftarrow Gen, b \leftarrow \{0, 1\}, c \leftarrow Enc_k(m_b)$
 3. $b' \leftarrow A(c)$Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Perfect indistinguishability

- Let $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Define a randomized experiment $\text{PrivK}_{A,\Pi}$:

1. A outputs $m_0, m_1 \in \mathcal{M}$
2. $k \leftarrow \text{Gen}$, $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Perfect indistinguishability

- Let $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Define a randomized experiment $\text{PrivK}_{A,\Pi}$:

1. A outputs $m_0, m_1 \in \mathcal{M}$
2. $k \leftarrow \text{Gen}, b \leftarrow \{0, 1\}, c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Π is *perfectly indistinguishable* if for *all* attackers (algorithms) A , it holds that

$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2$$

Perfect indistinguishability

- Let $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ be an encryption scheme with message space \mathcal{M} , and A an adversary

Define a randomized experiment $\text{PrivK}_{A,\Pi}$:

1. A outputs $m_0, m_1 \in \mathcal{M}$
2. $k \leftarrow \text{Gen}, b \leftarrow \{0, 1\}, c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Π is *perfectly indistinguishable* if for **all** attackers (algorithms) A , it holds that

$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2$$

Claim: Π is *perfectly indistinguishable* \Leftrightarrow Π is *perfectly secure*



Computational security?

- **Idea:** relax *perfect indistinguishability*



Computational security?

- **Idea:** relax *perfect indistinguishability*

Two approaches

- *Concrete* security
- *Asymptotic* security

Computational indistinguishability

- (t, ϵ) -indistinguishability (concrete)
 - Security may fail with probability $\leq \epsilon$
 - Restrict attention to attackers running in time $\leq t$



Computational indistinguishability

- (t, ϵ) -indistinguishability (concrete)
 - Security may fail with probability $\leq \epsilon$
 - Restrict attention to attackers running in time $\leq t$
- Π is (t, ϵ) -indistinguishable if for **all** attackers A running in time **at most** t , it holds that
$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2 + \epsilon$$



Computational indistinguishability

- (t, ϵ) -indistinguishability (concrete)
 - Security may fail with probability $\leq \epsilon$
 - Restrict attention to attackers running in time $\leq t$
- Π is (t, ϵ) -indistinguishable if for **all** attackers A running in time **at most** t , it holds that
$$\Pr[\text{PrivK}_{A,\Pi} = 1] \leq 1/2 + \epsilon$$

Does **not** lead to a clean theory ...

- Sensitive to exact computational model
- Π can be (t, ϵ) -secure for many choices of t, ϵ



Computational indistinguishability

- Introduce *security parameter* n (asymptotic)
 - For now, can view n as the *key length*
 - Fixed by honest parties at initialization
 - Known by adversary



Computational indistinguishability

- Introduce *security parameter* n (asymptotic)
 - For now, can view n as the *key length*
 - Fixed by honest parties at initialization
 - Known by adversary

Measure running time of all parties, and the *success probability* of the adversary, as *functions of n*

Computational indistinguishability

- Introduce *security parameter* n (asymptotic)
 - For now, can view n as the *key length*
 - Fixed by honest parties at initialization
 - Known by adversary

Measure running time of all parties, and the *success probability* of the adversary, as *functions of n*

Computational indistinguishability:

- Security may fail with probability *negligible* in n
- Restrict attention to attackers running in time (at most) *polynomial in n*



Definitions

- A function $f : \mathbb{Z}^+ \rightarrow \mathbb{Z}^+$ is (at most) *polynomial* if **there exists** c s.t. $f(n) < n^c$ for large enough n .

A function $f : \mathbb{Z}^+ \rightarrow [0, 1]$ is *negligible* if **every** polynomial p it holds that $f(n) < 1/p(n)$ for large enough n .

– Typical example: $f(n) = \text{poly}(n) \cdot 2^{-cn}$

Definitions

- A function $f : \mathbb{Z}^+ \rightarrow \mathbb{Z}^+$ is (at most) *polynomial* if **there exists** c s.t. $f(n) < n^c$ for large enough n .

A function $f : \mathbb{Z}^+ \rightarrow [0, 1]$ is *negligible* if **every** polynomial p it holds that $f(n) < 1/p(n)$ for large enough n .

– Typical example: $f(n) = \text{poly}(n) \cdot 2^{-cn}$

- “**Efficient**” = “(probabilistic) polynomial-time (**PPT**)” borrowed from *complexity theory*



Definitions

- A function $f : \mathbb{Z}^+ \rightarrow \mathbb{Z}^+$ is (at most) *polynomial* if **there exists** c s.t. $f(n) < n^c$ for large enough n .

A function $f : \mathbb{Z}^+ \rightarrow [0, 1]$ is *negligible* if **every** polynomial p it holds that $f(n) < 1/p(n)$ for large enough n .

– Typical example: $f(n) = \text{poly}(n) \cdot 2^{-cn}$

- “**Efficient**” = “(probabilistic) polynomial-time (**PPT**)” borrowed from *complexity theory*

- Convenient closure properties

– $\text{poly} * \text{poly} = \text{poly}$

– Poly-many calls to PPT subroutine (with poly-size input) is still PPT

– $\text{poly} * \text{negl} = \text{negl}$

– Poly-many calls to subroutine that fails with *negligible* probability fails with *negligible* probability overall



(Re)defining encryption

- A *private-key encryption scheme* is defined by three PPT algorithms (Gen, Enc, Dec):
 - Gen: takes as input 1^n ; outputs k
 - Enc: takes as input a key k and message $m \in \{0, 1\}^*$; outputs ciphertext c : $c \leftarrow \text{Enc}_k(m)$
 - Dec: takes key k and ciphertext c as input; outputs a message m or “error” (\perp)



Computational indistinguishability (asymptotic)

- Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}(n)$:

1. $A(1^n)$ outputs $m_0, m_1 \in \{0, 1\}^*$ of equal length
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.



Computational indistinguishability (asymptotic)

■ Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}(n)$:

1. $A(1^n)$ outputs $m_0, m_1 \in \{0, 1\}^*$ of equal length
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for *all PPT* attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$



Example

- Consider a scheme where the **best** attack is *brute-force search* over the key space, and $Gen(1^n)$ generates a uniform n -bit key
 - So if A runs in time $t(n)$, then
$$\Pr[PrivK_{A,\Pi}(n) = 1] < 1/2 + O(t(n)/2^n)$$



Example

- Consider a scheme where the **best** attack is *brute-force search* over the key space, and $Gen(1^n)$ generates a uniform n -bit key
 - So if A runs in time $t(n)$, then
$$\Pr[PrivK_{A,\Pi}(n) = 1] < 1/2 + O(t(n)/2^n)$$
 - The scheme is **EAV-secure**: for any polynomial t , the function $t(n)/2^n$ is **negligible**.



Example

- Consider a scheme and a particular attacker A that runs for n^3 minutes and breaks the scheme with probability $2^{40}2^{-n}$
 - This does **not** contradict asymptotic security



Example

- Consider a scheme and a particular attacker A that runs for n^3 minutes and breaks the scheme with probability $2^{40}2^{-n}$
 - This does **not** contradict asymptotic security
 - What about real-world security (against this attacker)?
 - $n = 40$: A breaks with prob. 1 in 6 weeks
 - $n = 50$: A breaks with prob. 1/1000 in 3 months
 - $n = 500$: A breaks with prob. 2^{-500} in 200 years



Example

- What happens when computers get faster?
 - Consider a scheme that takes time n^2 to run but time 2^n to break with prob. 1



Example

- What happens when computers get faster?
 - Consider a scheme that takes time n^2 to run but time 2^n to break with prob. 1

What if computers get $4\times$ faster?

- Users double n ; maintain same running time
- Attacker's work is (roughly) squared!



Encryption and plaintext length

- In practice, we want encryption schemes that can encrypt **arbitrary-length** messages.



Encryption and plaintext length

- In practice, we want encryption schemes that can encrypt **arbitrary-length** messages.
- In general, encryption does **not** hide the plaintext length
 - The definition takes this into account by requiring m_0 , m_1 to have the **same** length.



Encryption and plaintext length

- In practice, we want encryption schemes that can encrypt **arbitrary-length** messages.
- In general, encryption does **not** hide the plaintext length
 - The definition takes this into account by requiring m_0 , m_1 to have the **same** length.
- But leaking plaintext length can **often** lead to problems in the real world!
 - Databases searches
 - Encrypting compressed data



Micali & Goldwasser



Silvio Micali



Shafi Goldwasser

1984: semantic security, indistinguishability
(Turing Award 2012)

Micali & Blum



Silvio Micali



Manuel Blum

1984: defined notion of pseudo-random generator
(Turing Award 1995)

Pseudorandomness

- Important building block for computationally secure encryption



Pseudorandomness

- Important building block for computationally secure encryption
- What does “random” mean?



Pseudorandomness

- Important **building block** for **computationally secure encryption**
- What does “**random**” mean?
- Which of the following is a **uniform** string?
 - 0101010101010101
 - 0010111011100110
 - 0000000000000000



Pseudorandomness

- Important building block for computationally secure encryption
- What does “random” mean?
- Which of the following is a uniform string?
 - 0101010101010101
 - 0010111011100110
 - 0000000000000000
- If we generate a uniform 16-bit string, each of the above occurs with probability 2^{-16}



What does “pseudorandom” mean?

- Informal: **Cannot** be distinguished from **uniform** (“random”)



What does “pseudorandom” mean?

- Informal: **Cannot** be distinguished from **uniform** (“random”)
- Which of the following is **pseudorandom**?
 - 0101010101010101
 - 0010111011100110
 - 0000000000000000

What does “pseudorandom” mean?

- Informal: **Cannot** be distinguished from **uniform** (“random”)
- Which of the following is **pseudorandom**?
 - 0101010101010101
 - 0010111011100110
 - 0000000000000000
- *Pseudorandomness* is a property of a *distribution*, **not** a string.



Pseudorandomness

- Fix some distribution D on n -bit strings
 - $x \leftarrow D$ means “sample x according to D ”



Pseudorandomness

- Fix some distribution D on n -bit strings
 - $x \leftarrow D$ means “sample x according to D ”
- Historically, D was considered *pseudorandom* if it “passed a bunch of statistical tests”
 - $\Pr_{x \leftarrow D}[1^{st} \text{ bit of } x \text{ is } 1] \approx 1/2$
 - $\Pr_{x \leftarrow D}[\text{parity of } x \text{ is } 1] \approx 1/2$
 - $\Pr_{x \leftarrow D}[A_i(x) = 1] \approx \Pr_{x \leftarrow U_n}[A_i(x) = 1]$
for $i = 1, \dots, 20$



Pseudorandomness

- Fix some distribution D on n -bit strings
 - $x \leftarrow D$ means “sample x according to D ”
- Historically, D was considered *pseudorandom* if it “passed a bunch of statistical tests”
 - $\Pr_{x \leftarrow D}[1^{\text{st}} \text{ bit of } x \text{ is } 1] \approx 1/2$
 - $\Pr_{x \leftarrow D}[\text{parity of } x \text{ is } 1] \approx 1/2$
 - $\Pr_{x \leftarrow D}[A_i(x) = 1] \approx \Pr_{x \leftarrow U_n}[A_i(x) = 1]$
for $i = 1, \dots, 20$

This is **not** sufficient, since it is **not** possible to know what statistical test an attacker will use.



Pseudorandomness

- Cryptographic definition of *pseudorandomness*
 - D is *pseudorandom* if it passes all *efficient* statistical tests



Pseudorandomness

- Cryptographic definition of *pseudorandomness*
 - D is *pseudorandom* if it passes all *efficient* statistical tests

(Concrete) Let D be a distribution on p -bit strings. D is (t, ϵ) -*pseudorandom* if for all A running in time **at most** t ,

$$|\Pr_{x \leftarrow D}[A(x) = 1] - \Pr_{x \leftarrow U_p}[A(x) = 1]| \leq \epsilon$$



Pseudorandomness

- Cryptographic definition of *pseudorandomness*
 - D is *pseudorandom* if it passes all *efficient* statistical tests

(Concrete) Let D be a distribution on p -bit strings. D is *(t, ϵ) -pseudorandom* if for all A running in time **at most** t ,

$$|\Pr_{x \leftarrow D}[A(x) = 1] - \Pr_{x \leftarrow U_p}[A(x) = 1]| \leq \epsilon$$

(Asymptotic) *Security parameter* n , polynomial p

Definiton 3.2 Let D_n be a distribution over $p(n)$ -bit strings. $\{D_n\}$ is *pseudorandom* if for all probabilistic, polynomial-time (PPT) distinguishers A , there is a **negligible** function ϵ such that

$$|\Pr_{x \leftarrow D_n}[A(x) = 1] - \Pr_{x \leftarrow U_{p(n)}}[A(x) = 1]| \leq \epsilon(n)$$

Pseudorandom generators (PRGs)

- A *PRG* is an *efficient*, *deterministic* algorithm that expands a *short, uniform seed* into a *longer, pseudorandom* output



Pseudorandom generators (PRGs)

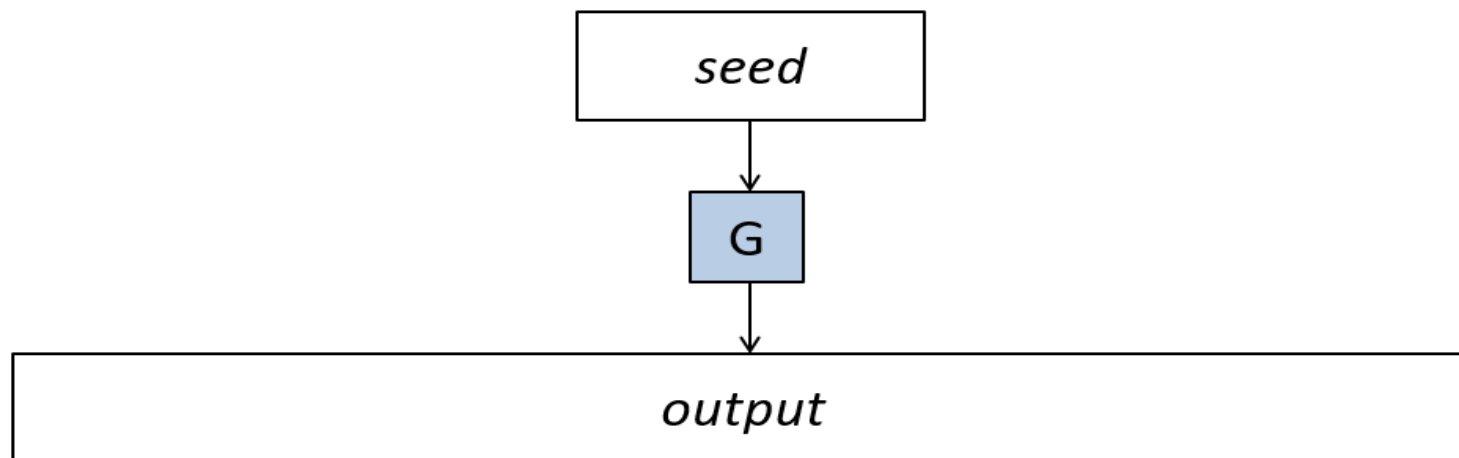
- A *PRG* is an *efficient*, *deterministic* algorithm that expands a *short, uniform seed* into a *longer, pseudorandom* output
 - Useful whenever you have a “*small*” number of true random bits, and want lots of “*random-looking*” bits



Pseudorandom generators (PRGs)

- A *PRG* is an *efficient*, *deterministic* algorithm that expands a *short, uniform seed* into a *longer, pseudorandom* output
 - Useful whenever you have a “*small*” number of true random bits, and want lots of “*random-looking*” bits

Let G be a deterministic, poly-time algorithm that is *expanding*, i.e., $|G(x)| = p(|x|) > |x|$.



Pseudorandom generators (PRGs)

- A *PRG* is an **efficient**, **deterministic** algorithm that expands a *short, uniform seed* into a *longer, pseudorandom* output
 - Useful whenever you have a “**small**” number of true random bits, and want lots of “**random-looking**” bits

Let G be a deterministic, poly-time algorithm that is *expanding*, i.e., $|G(x)| = p(|x|) > |x|$.

G defines a sequence of distributions.

- D_n = the distribution on $p(n)$ -bit strings defined by choosing $x \leftarrow U_n$ and outputting $G(x)$
- $\Pr_{D_n}[y] = \Pr_{U_n}[G(x) = y] = \sum_{x: G(x)=y} \Pr_{U_n}[x]$
$$= \sum_{x: G(x)=y} 2^{-n}$$
$$= |\{x : G(x) = y\}| / 2^n$$



- For all **efficient** distinguishers A , there is a **negligible** function ϵ such that

$$|\Pr_{x \leftarrow U_n}[A(G(x)) = 1] - \Pr_{y \leftarrow U_{p(n)}}[A(y) = 1]| \leq \epsilon(n)$$

- For all **efficient** distinguishers A , there is a **negligible** function ϵ such that

$$|\Pr_{x \leftarrow U_n}[A(G(x)) = 1] - \Pr_{y \leftarrow U_{p(n)}}[A(y) = 1]| \leq \epsilon(n)$$

No efficient A can distinguish whether it is given $G(x)$ (for uniform x) or a uniform string y !

- For all **efficient** distinguishers A , there is a **negligible** function ϵ such that

$$|\Pr_{x \leftarrow U_n}[A(G(x)) = 1] - \Pr_{y \leftarrow U_{p(n)}}[A(y) = 1]| \leq \epsilon(n)$$

No efficient A can distinguish whether it is given $G(x)$ (for uniform x) or a uniform string y !

- PRGs are **limited**
 - They have **fixed-length** output
 - They produce the entire output in “**one shot**”
 - In practice, PRGs are based on **stream ciphers**
 - Can be viewed as producing an “**unbounded**” stream of pseudorandom bits, on demand
 - Will be revisited later

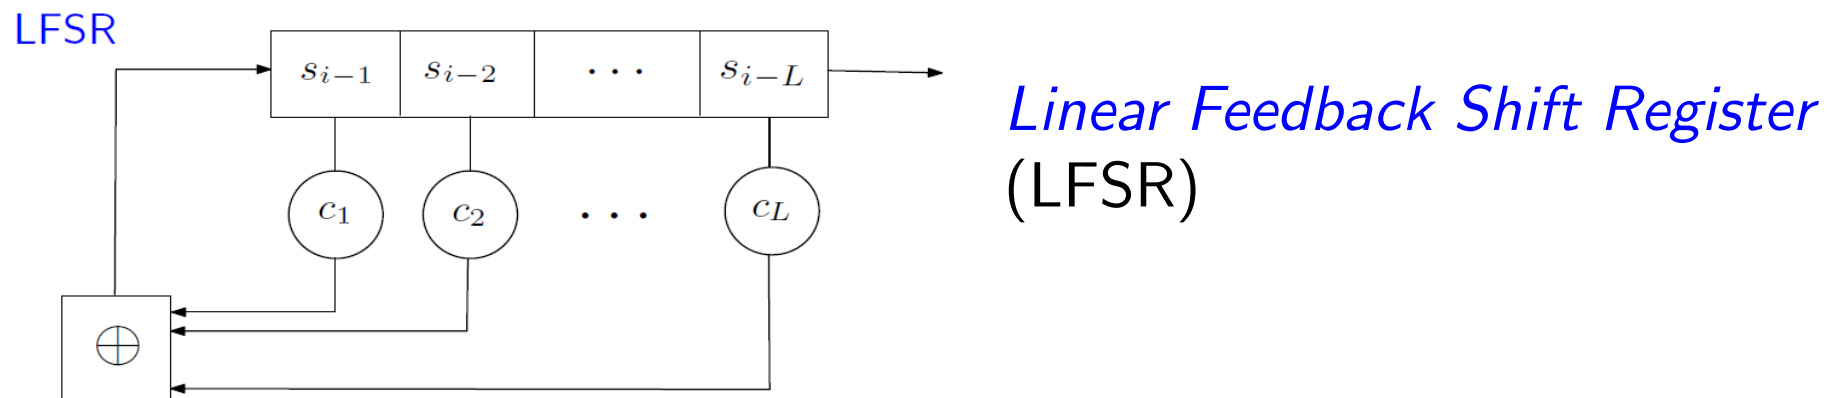
Do PRGs/stream ciphers exist?

- We **don't** know ...
 - Would imply $P \neq NP$
 - We will **assume** certain algorithms are PRGs
 - Can construct PRGs from **weaker** assumptions (later)




Do PRGs/stream ciphers exist?

- We **don't** know ...
 - Would imply $P \neq NP$
 - We will **assume** certain algorithms are PRGs
 - Can construct PRGs from **weaker** assumptions (later)



Example

$s = (000100110101111)^{15}$




$$LC(s) = 4$$

$$P_s(x) = x^4 + x + 1$$

Practical “PRGs”

■ RC4

```
i := 0
j := 0
while GeneratingOutput:
  i := (i + 1) mod 256
  j := (j + S[i]) mod 256
  swap values of S[i] and S[j]
  K := S[(S[i] + S[j]) mod 256]
  output K
endwhile
```



Blum-Blum-Shub

```
num_outputted = 0;
while num_outputted < m:
  X := X*X mod N
  num_outputted := num_outputted + 1
  output least-significant-bit(X)
```

SIAM J. COMPUT.
Vol. 15, No. 2, May 1986

© 1986 Society for Industrial and Applied Mathematics
003

A SIMPLE UNPREDICTABLE PSEUDO-RANDOM NUMBER GENERATOR*

L. BLUM†, M. BLUM‡ AND M. SHUB§

Abstract. Two closely-related pseudo-random sequence generators are presented: The $1/P$ generator, with input P a prime, outputs the quotient digits obtained on dividing 1 by P . The $x^2 \bmod N$ generator with inputs N, x_0 (where $N = P \cdot Q$ is a product of distinct primes, each congruent to 3 mod 4, and x_0 is a quadratic residue mod N), outputs $b_0 b_1 b_2 \dots$ where b_i = parity (x_i) and $x_{i+1} = x_i^2 \bmod N$.

From short seeds each generator efficiently produces long well-distributed sequences. Moreover, both generators have computationally hard problems at their core. The first generator's sequences, however, are *completely predictable* (from any small segment of $2|P|+1$ consecutive digits one can infer the “seed,” P , and continue the sequence backwards and forwards), whereas the second, under a certain intractability assumption, is *unpredictable* in a precise sense. The second generator has additional interesting properties: from knowledge of x_0 and N but *not* P or Q , one can generate the sequence forwards, but, under the above-mentioned intractability assumption, one can *not* generate the sequence backwards. From the additional knowledge of P and Q , one *can* generate the sequence backwards; one can even “jump” about from any point in the sequence to any other. Because of these properties, the $x^2 \bmod N$ generator promises many interesting applications, e.g., to public-key cryptography. To use these generators in practice, an analysis is needed of various properties of these sequences such as their periods. This analysis is begun here.

Key words. random, pseudo-random, Monte Carlo, computational complexity, secure transactions, public-key encryption, cryptography, one-time pad, Jacobi symbol, quadratic residuacity

Where things stand

- We saw that there are some inherent **limitations** if we want *perfect security*
 - In particular, key must be as **long** as the message



Where things stand


- We saw that there are some inherent **limitations** if we want *perfect security*
 - In particular, key must be as **long** as the message

We defined *computational security*, a **relaxed** notion of security



Q: Can we overcome prior limitations?



Recall: one-time pad



Recall: one-time pad



Pseudo one-time pad

- Let G be a deterministic, with $|G(k)| = p(|k|)$

$Gen(1^n)$: output uniform n -bit key k

– Security parameter $n \Rightarrow$ message space $\{0, 1\}^{p(n)}$

$Enc_k(m)$: output $G(k) \oplus m$

$Dec_k(m)$: output $G(k) \oplus c$



Pseudo one-time pad

- Let G be a deterministic, with $|G(k)| = p(|k|)$

$Gen(1^n)$: output uniform n -bit key k

- Security parameter $n \Rightarrow$ message space $\{0, 1\}^{p(n)}$

$Enc_k(m)$: output $G(k) \oplus m$

$Dec_k(m)$: output $G(k) \oplus c$

- Would like to be able to prove *computational security*
 - Based on the **assumption** that G is a PRG



Proof by reduction




Figure 2.1: We show that the security of S' implies the security of S by transforming an adversary Eve breaking S into an adversary Eve' breaking S'

Eve breaks $S \rightarrow Eve'$ breaks S'
 S' is secure $\rightarrow S$ is secure

Proof by reduction

- 1. Assume that G is a *PRG*
- 2. Assume toward a **contradiction** that there is an **efficient attacker** A who “breaks” the pseudo-OTP scheme
- 3. Use A as a **subroutine** to build an efficient D that “breaks” *pseudorandomness* of G



Proof by reduction

- 1. Assume that G is a *PRG*
 - 2. Assume toward a **contradiction** that there is an **efficient attacker** A who “breaks” the pseudo-OTP scheme
 - 3. Use A as a **subroutine** to build an efficient D that “breaks” *pseudorandomness* of G
 - By assumption, **no such** D exists!
- ⇒ **No** such A can exist



Proof by reduction

- 1. Assume that G is a *PRG*
 - 2. Assume toward a **contradiction** that there is an **efficient attacker** A who “breaks” the pseudo-OTP scheme
 - 3. Use A as a **subroutine** to build an efficient D that “breaks” *pseudorandomness* of G
 - By assumption, **no such D exists!**
- ⇒ **No** such A can exist

Theorem 3.3 If G is a pseudorandom generator (PRG), then the pseudo one-time pad (pseudo-OTP) Π is *EAV-secure* (i.e., *computationally secure*)



Next Lecture

- Pseudorandom functions, block ciphers ...

