



CSE5014 CRYPTOGRAPHY AND NETWORK SECURITY

Dr. QI WANG

Department of Computer Science and Engineering

Office: Room413, CoE South Tower

Email: wangqi@sustech.edu.cn

Pseudorandomness

- Cryptographic definition of *pseudorandomness*
 - D is *pseudorandom* if it passes all *efficient* statistical tests

Pseudorandomness

- Cryptographic definition of *pseudorandomness*
 - D is *pseudorandom* if it passes all *efficient* statistical tests

(Concrete) Let D be a distribution on p -bit strings. D is (t, ϵ) -*pseudorandom* if for all A running in time **at most** t ,

$$|\Pr_{x \leftarrow D}[A(x) = 1] - \Pr_{x \leftarrow U_p}[A(x) = 1]| \leq \epsilon$$



Pseudorandomness

- Cryptographic definition of *pseudorandomness*
 - D is *pseudorandom* if it passes all *efficient* statistical tests

(Concrete) Let D be a distribution on p -bit strings. D is (t, ϵ) -*pseudorandom* if for all A running in time **at most** t ,

$$|\Pr_{x \leftarrow D}[A(x) = 1] - \Pr_{x \leftarrow U_p}[A(x) = 1]| \leq \epsilon$$

(Asymptotic) *Security parameter* n , polynomial p

Definiton 3.2 Let D_n be a distribution over $p(n)$ -bit strings. $\{D_n\}$ is *pseudorandom* if for all probabilistic, polynomial-time (PPT) distinguishers A , there is a **negligible** function ϵ such that

$$|\Pr_{x \leftarrow D_n}[A(x) = 1] - \Pr_{x \leftarrow U_{p(n)}}[A(x) = 1]| \leq \epsilon(n)$$

- A *PRG* is an *efficient*, *deterministic* algorithm that expands a *short, uniform seed* into a *longer, pseudorandom* output

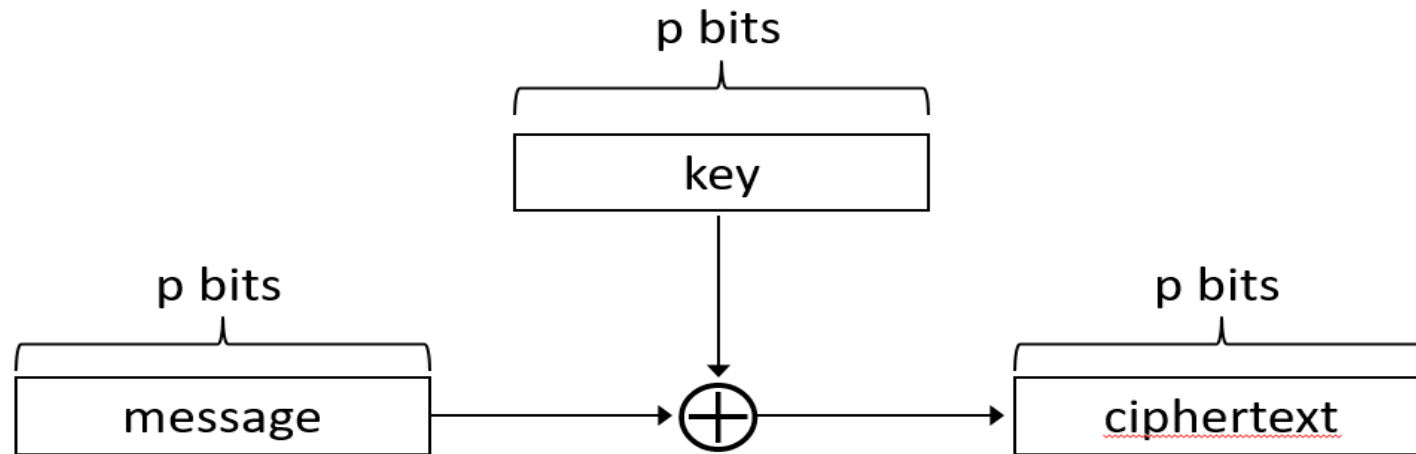
Let G be a deterministic, poly-time algorithm that is *expanding*, i.e., $|G(x)| = p(|x|) > |x|$.

- For all *efficient* distinguishers A , there is a *negligible* function ϵ such that

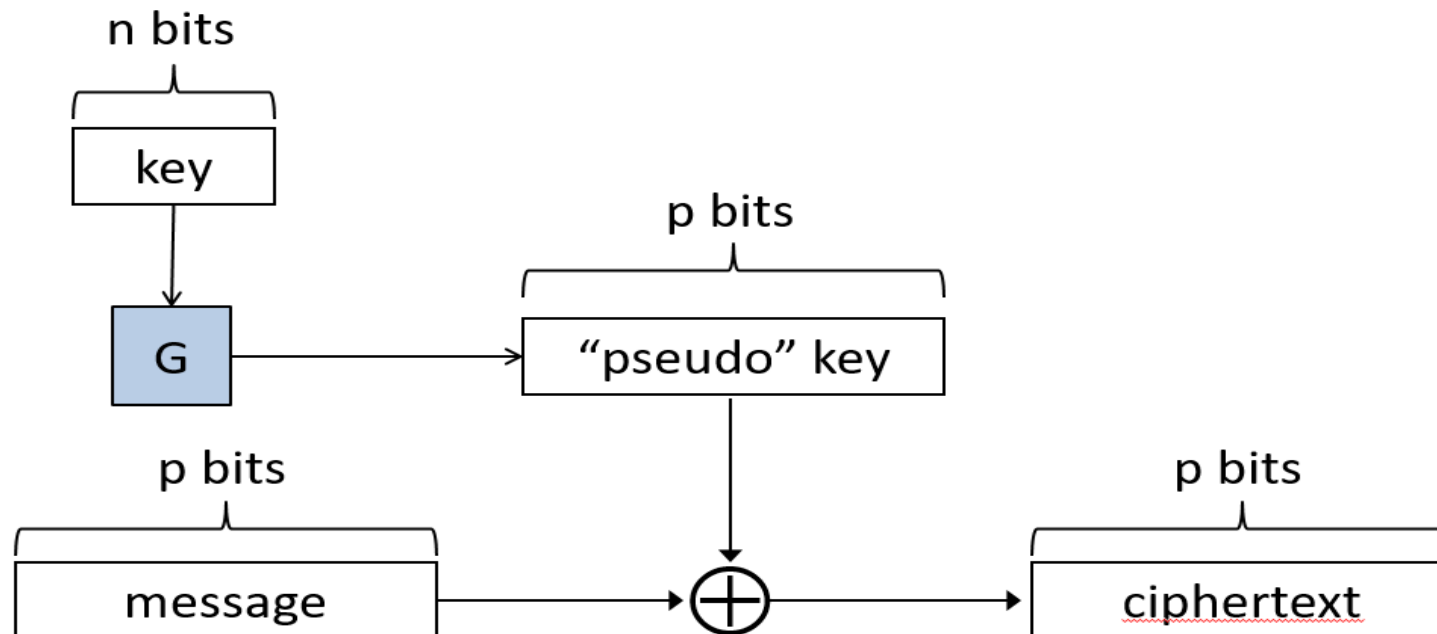
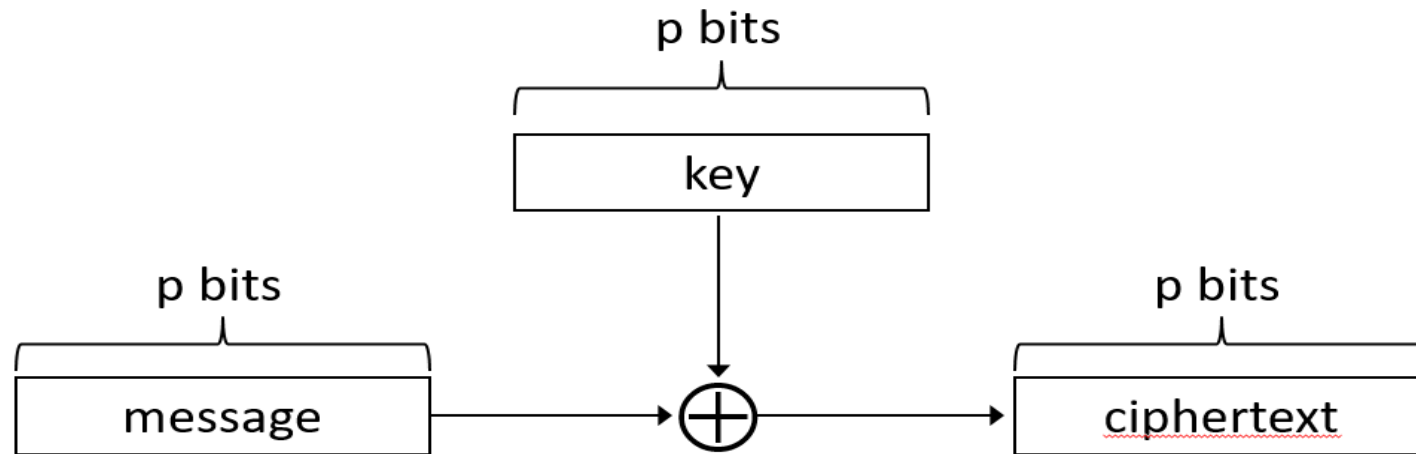
$$|\Pr_{x \leftarrow U_n}[A(G(x)) = 1] - \Pr_{y \leftarrow U_{p(n)}}[A(y) = 1]| \leq \epsilon(n)$$

No efficient A can distinguish whether it is given $G(x)$ (for uniform x) or a uniform string y !

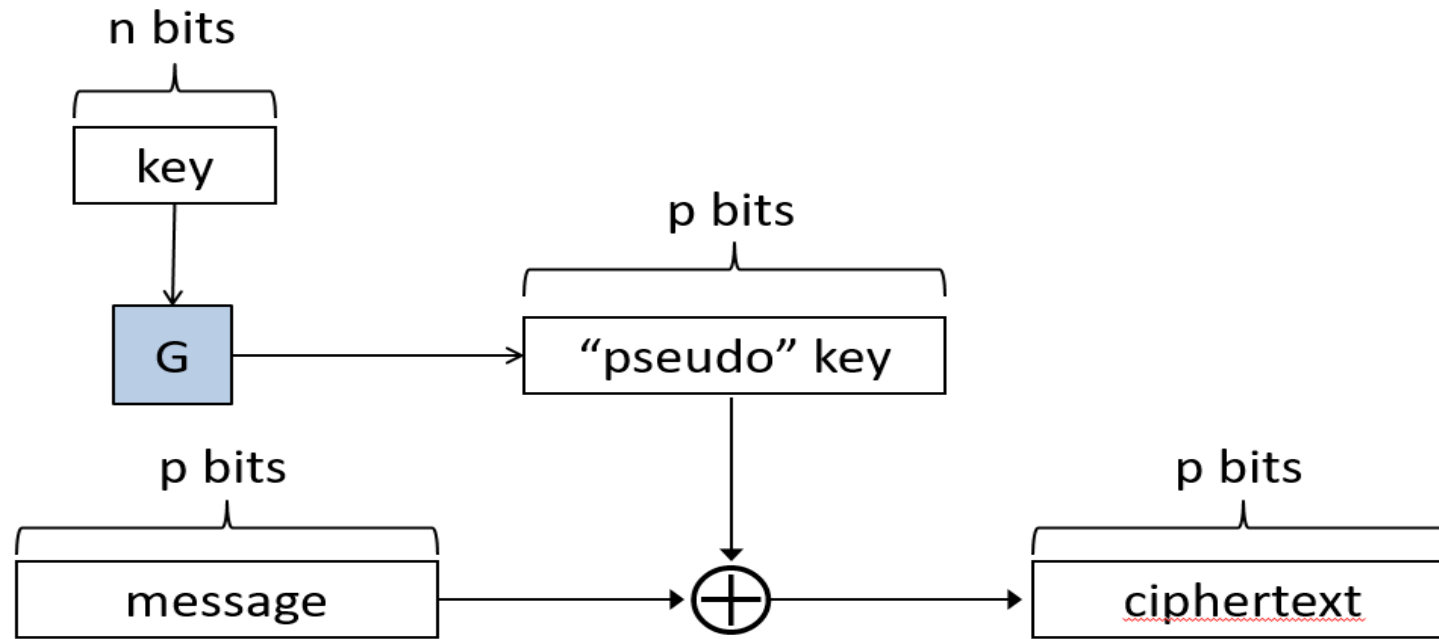
Recall: one-time pad



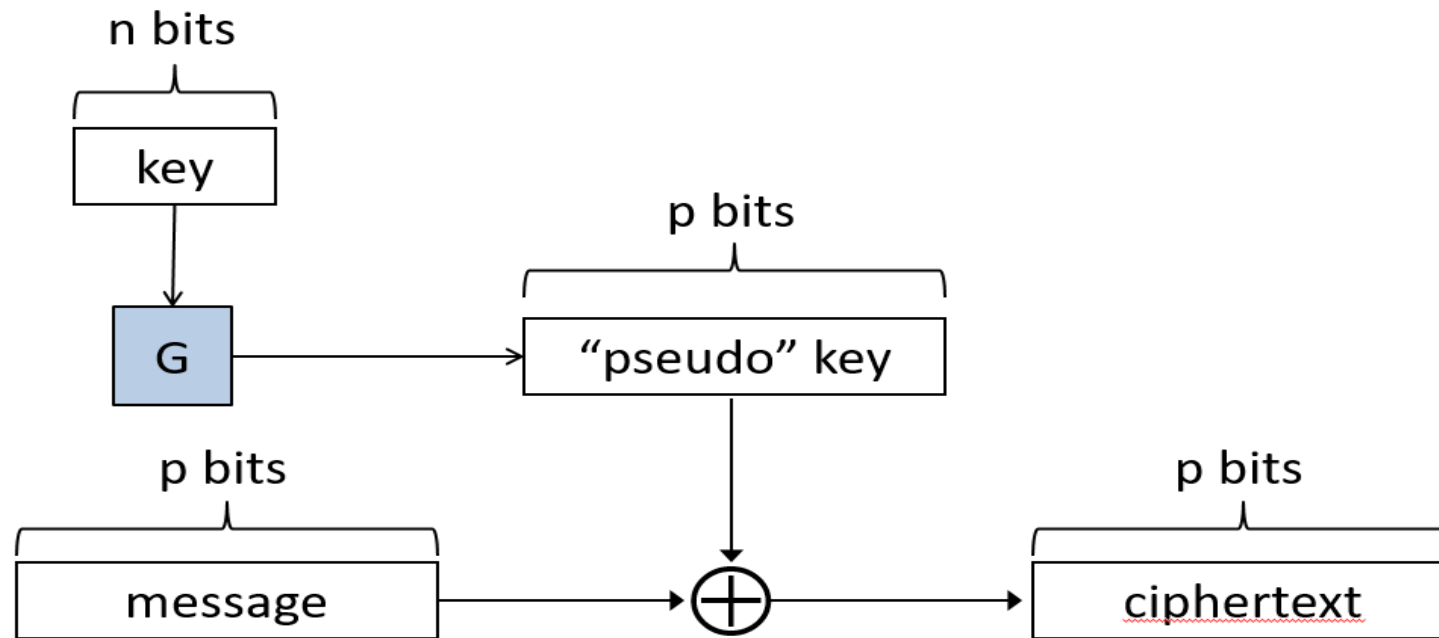
Recall: one-time pad



Pseudo one-time pad



Pseudo one-time pad



- Let G be a deterministic, with $|G(k)| = p(|k|)$
 $Gen(1^n)$: output uniform n -bit key k
 - Security parameter $n \Rightarrow$ message space $\{0, 1\}^{p(n)}$ $Enc_k(m)$: output $G(k) \oplus m$
 $Dec_k(m)$: output $G(k) \oplus c$

Proof by reduction

- 1. Assume that G is a *PRG*
 - 2. Assume toward a **contradiction** that there is an **efficient attacker** A who “breaks” the pseudo-OTP scheme
 - 3. Use A as a **subroutine** to build an efficient D that “breaks” *pseudorandomness* of G
 - By assumption, **no such** D exists!
- ⇒ **No** such A can exist

Proof by reduction

- 1. Assume that G is a *PRG*
 - 2. Assume toward a **contradiction** that there is an **efficient attacker** A who “breaks” the pseudo-OTP scheme
 - 3. Use A as a **subroutine** to build an efficient D that “breaks” *pseudorandomness* of G
 - By assumption, **no such D exists!**
- ⇒ **No** such A can exist

Theorem 3.3 If G is a pseudorandom generator (PRG), then the pseudo one-time pad (pseudo-OTP) Π is *EAV-secure* (i.e., *computationally secure*)



The reduction

- **Proof.**

The reduction

■ Proof.

Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}(n)$:

1. $A(1^n)$ outputs $m_0, m_1 \in \{0, 1\}^*$ of equal length
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$
3. $b' \leftarrow A(c)$

Adversary A *succeeds* if $b = b'$, and we say the experiment evaluates to 1 in this case.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for *all PPT* attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$



The reduction

■ Proof.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for **all PPT** attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$

The reduction

■ Proof.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for **all PPT** attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$

Suppose to the **contrary** that there exists an **efficient** attacker A , s.t.

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] > 1/2 + 1/\text{poly}(n)$$

The reduction

■ Proof.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for **all PPT** attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$

Suppose to the **contrary** that there exists an **efficient** attacker A , s.t.

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] > 1/2 + 1/\text{poly}(n)$$

This means, $\Pr[A(\text{Enc}_{U_n}(m)) = 1] - \Pr[A(U_{p(n)}) = 1] > 1/\text{poly}(n)$

The reduction

■ Proof.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for **all PPT** attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$

Suppose to the **contrary** that there exists an **efficient** attacker A , s.t.

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] > 1/2 + 1/\text{poly}(n)$$

This means, $\Pr[A(\text{Enc}_{U_n}(m)) = 1] - \Pr[A(U_{p(n)}) = 1] > 1/\text{poly}(n)$

$$|\Pr[A(G(U_n) \oplus m) = 1] - \Pr[A(U_{p(n)}) = 1]| > 1/\text{poly}(n)$$

The reduction

■ Proof.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for **all PPT** attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$

Suppose to the **contrary** that there exists an **efficient** attacker A , s.t.

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] > 1/2 + 1/\text{poly}(n)$$

This means, $\Pr[A(\text{Enc}_{U_n}(m)) = 1] - \Pr[A(U_{p(n)}) = 1] > 1/\text{poly}(n)$

$$|\Pr[A(G(U_n) \oplus m) = 1] - \Pr[A(U_{p(n)}) = 1]| > 1/\text{poly}(n)$$

$\text{Enc}_{U_n}()$ is **not pseudorandom**

The reduction

■ Proof.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for **all PPT** attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$

Suppose to the **contrary** that there exists an **efficient** attacker A , s.t.

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] > 1/2 + 1/\text{poly}(n)$$

This means, $\Pr[A(\text{Enc}_{U_n}(m)) = 1] - \Pr[A(U_{p(n)}) = 1] > 1/\text{poly}(n)$

$$|\Pr[A(G(U_n) \oplus m) = 1] - \Pr[A(U_{p(n)}) = 1]| > 1/\text{poly}(n)$$

$\text{Enc}_{U_n}()$ is **not pseudorandom**

Define $D : \{0, 1\}^{p(n)} \rightarrow \{0, 1\}$ as: $D(y) = A(y \oplus m)$, which means $A(z) = D(z \oplus m)$. Note that D is also **efficient**. But we have

$$|\Pr[D(G(U_n)) = 1] - \Pr[D(U_{p(n)} \oplus m) = 1]| \geq 1/\text{poly}(n)$$

The reduction

■ Proof.

Definition 3.1 Π is *computationally indistinguishable* (aka *EAV-secure*) if for **all PPT** attackers (algorithms) A , there is a *negligible* function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$

Suppose to the **contrary** that there exists an **efficient** attacker A , s.t.

$$\Pr[\text{PrivK}_{A,\Pi}(n) = 1] > 1/2 + 1/\text{poly}(n)$$

This means, $\Pr[A(\text{Enc}_{U_n}(m)) = 1] - \Pr[A(U_{p(n)}) = 1] > 1/\text{poly}(n)$

$$|\Pr[A(G(U_n) \oplus m) = 1] - \Pr[A(U_{p(n)}) = 1]| > 1/\text{poly}(n)$$

$\text{Enc}_{U_n}()$ is **not pseudorandom**

Define $D : \{0, 1\}^{p(n)} \rightarrow \{0, 1\}$ as: $D(y) = A(y \oplus m)$, which means $A(z) = D(z \oplus m)$. Note that D is also **efficient**. But we have

$$|\Pr[D(G(U_n)) = 1] - \Pr[D(U_{p(n)} \oplus m) = 1]| \geq 1/\text{poly}(n)$$

Since $U_{p(n)} \oplus m \equiv U_{p(n)}$, this **contradicts** that G is a PRG.

Proof by reduction (alternatively)

- 1. Assume that G is a *PRG*
- 2. Fix some *arbitrary*, efficient A attacking the pseudo-OTP scheme
- 3. Use A as a *subroutine* to build an efficient D attacking G
 - Relate the distinguishing probability of D to the success probability of A
- 4. By assumption, the distinguishing probability of D *must be negligible*



Proof by reduction (alternatively)

- 1. Assume that G is a *PRG*
- 2. Fix some *arbitrary*, efficient A attacking the pseudo-OTP scheme
- 3. Use A as a *subroutine* to build an efficient D attacking G
 - Relate the distinguishing probability of D to the success probability of A
- 4. By assumption, the distinguishing probability of D *must be negligible*
 - \Rightarrow *Bound* the success probability of A



The reduction

- Let $\mu(n) = \Pr[\text{PrivK}_{A,\Pi}(n) = 1]$



The reduction

- Let $\mu(n) = \Pr[\text{PrivK}_{A,\Pi}(n) = 1]$

If distribution of y is **pseudorandom**, then the view of A is *exactly* the same as in $\text{PrivK}_{A,\Pi}(n)$

$$\Rightarrow \Pr_{x \leftarrow U_n}[D(G(x)) = 1] = \mu(n)$$

The reduction

- Let $\mu(n) = \Pr[\text{PrivK}_{A,\Pi}(n) = 1]$

If distribution of y is **pseudorandom**, then the view of A is *exactly* the same as in $\text{PrivK}_{A,\Pi}(n)$

$$\Rightarrow \Pr_{x \leftarrow U_n}[D(G(x)) = 1] = \mu(n)$$

If distribution of y is **uniform**, then A succeeds with probability exactly $1/2$

$$\Rightarrow \Pr_{y \leftarrow U_{p(n)}}[D(y) = 1] = 1/2$$

The reduction

- Let $\mu(n) = \Pr[\text{PrivK}_{A,\Pi}(n) = 1]$

If distribution of y is **pseudorandom**, then the view of A is *exactly* the same as in $\text{PrivK}_{A,\Pi}(n)$

$$\Rightarrow \Pr_{x \leftarrow U_n}[D(G(x)) = 1] = \mu(n)$$

If distribution of y is **uniform**, then A succeeds with probability exactly $1/2$

$$\Rightarrow \Pr_{y \leftarrow U_{p(n)}}[D(y) = 1] = 1/2$$

Since G is pseudorandom,

$$|\mu(n) - 1/2| \leq \text{negl}(n)$$

The reduction

- Let $\mu(n) = \Pr[\text{PrivK}_{A,\Pi}(n) = 1]$

If distribution of y is **pseudorandom**, then the view of A is *exactly* the same as in $\text{PrivK}_{A,\Pi}(n)$

$$\Rightarrow \Pr_{x \leftarrow U_n}[D(G(x)) = 1] = \mu(n)$$

If distribution of y is **uniform**, then A succeeds with probability exactly $1/2$

$$\Rightarrow \Pr_{y \leftarrow U_{p(n)}}[D(y) = 1] = 1/2$$

Since G is pseudorandom,

$$|\mu(n) - 1/2| \leq \text{negl}(n)$$

$$\Rightarrow \Pr[\text{PrivK}_{A,\Pi}(n) = 1] \leq 1/2 + \text{negl}(n)$$

Have we gained anything?

- YES: the pseudo-OTP has a key **shorter** than the message
 - n bits vs. $p(n)$ bits

Have we gained anything?

- YES: the pseudo-OTP has a key **shorter** than the message
 - n bits vs. $p(n)$ bits

Recall: Perfect security has two **limitations**

- Key as long as the message
- Key can only be used once



Have we gained anything?

- YES: the pseudo-OTP has a key **shorter** than the message
 - n bits vs. $p(n)$ bits

Recall: Perfect security has two **limitations**

- Key as long as the message
- Key can only be used once

The pseudo OTP **still has** the second limitation (for the same reason as the OTP)



Have we gained anything?

- YES: the pseudo-OTP has a key **shorter** than the message
 - n bits vs. $p(n)$ bits

Recall: Perfect security has two **limitations**

- Key as long as the message
- Key can only be used once

The pseudo OTP **still has** the second limitation (for the same reason as the OTP)

How can we circumvent the second limitation?



But first...

- Develop an appropriate security *definition*
 - Security goal
 - Threat model

But first...

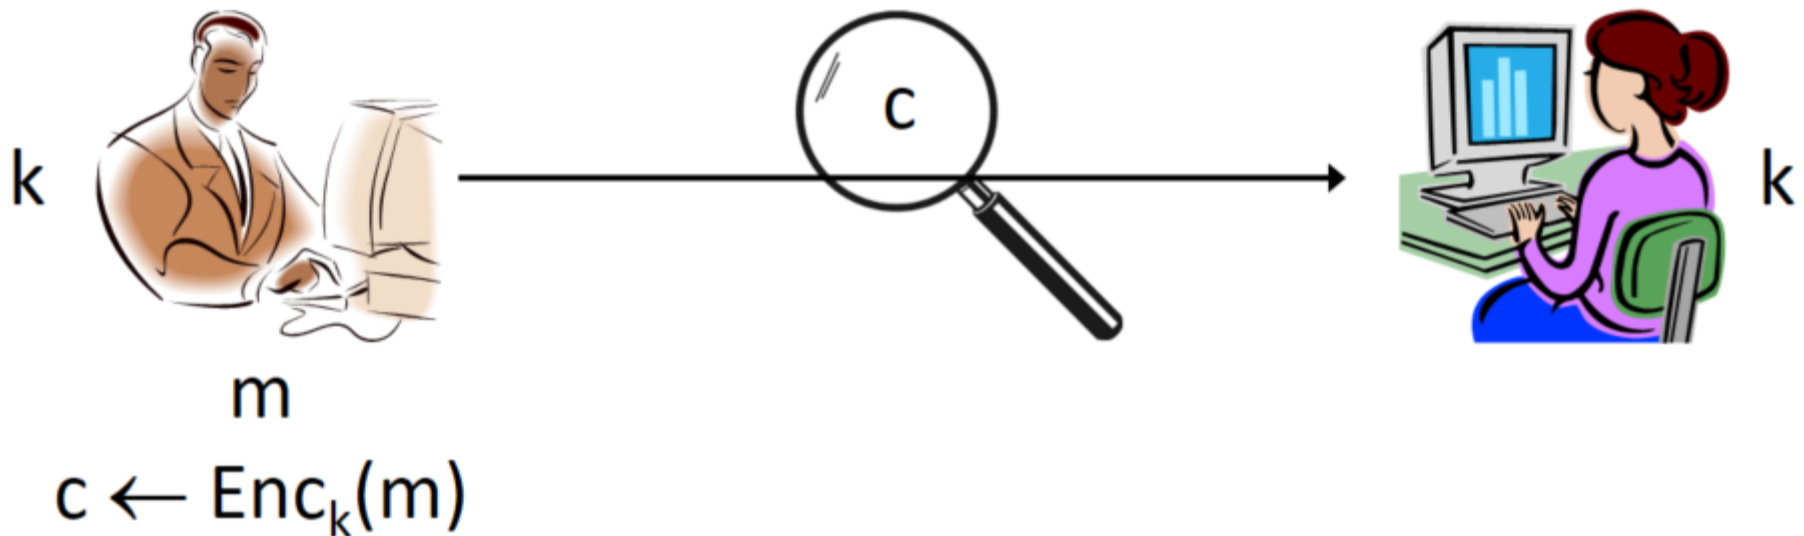
- Develop an appropriate security *definition*
 - Security goal
 - Threat model

We will keep the security goal the same, but **strengthen** the threat model

But first...

- Develop an appropriate security *definition*
 - Security goal
 - Threat model

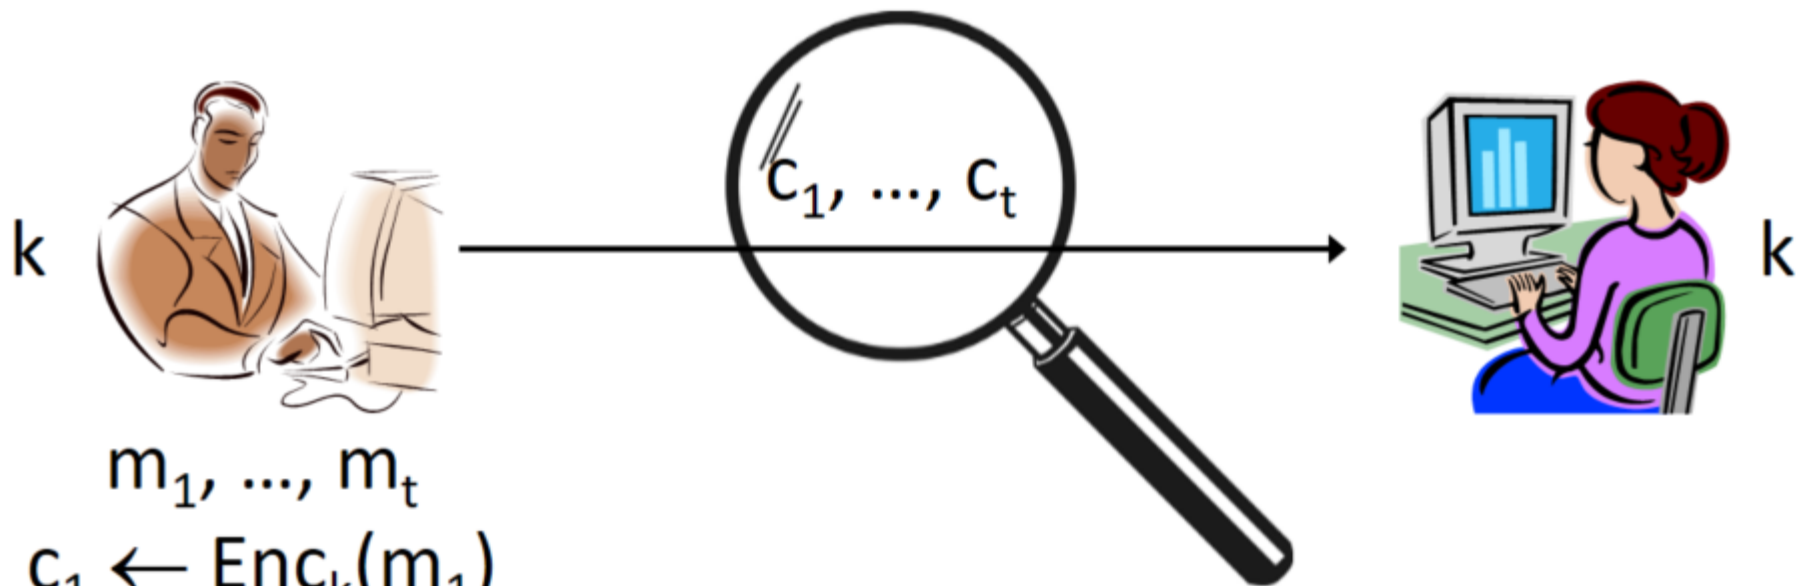
We will keep the security goal the same, but **strengthen** the threat model



But first...

- Develop an appropriate security *definition*
 - Security goal
 - Threat model

We will keep the security goal the same, but **strengthen** the threat model



A formal definition

- Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}^{\text{mult}}(n)$:

1. $A(1^n)$ outputs two **vectors** $(m_{0,1}, \dots, m_{0,t})$ and $(m_{1,1}, \dots, m_{1,t})$
Required that $|m_{0,i}| = |m_{1,i}|$ for all i
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, for all i , $c_i \leftarrow \text{Enc}_k(m_{b,i})$
3. $b' \leftarrow A(c_1, \dots, c_t)$

Adversary A **succeeds** if $b = b'$, and the experiment evaluates to **1** in this case.



A formal definition

- Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}^{\text{mult}}(n)$:

1. $A(1^n)$ outputs two **vectors** $(m_{0,1}, \dots, m_{0,t})$ and $(m_{1,1}, \dots, m_{1,t})$
Required that $|m_{0,i}| = |m_{1,i}|$ for all i
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, for all i , $c_i \leftarrow \text{Enc}_k(m_{b,i})$
3. $b' \leftarrow A(c_1, \dots, c_t)$

Adversary A **succeeds** if $b = b'$, and the experiment evaluates to **1** in this case.

Definition 3.4 Π is **multiple-message indistinguishable** if for **all PPT** attackers A , there is a **negligible** function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}^{\text{mult}}(n) = 1] \leq 1/2 + \epsilon(n)$$



A formal definition

- Fix Π , A

Define a randomized experiment $\text{PrivK}_{A,\Pi}^{\text{mult}}(n)$:

1. $A(1^n)$ outputs two **vectors** $(m_{0,1}, \dots, m_{0,t})$ and $(m_{1,1}, \dots, m_{1,t})$
Required that $|m_{0,i}| = |m_{1,i}|$ for all i
2. $k \leftarrow \text{Gen}(1^n)$, $b \leftarrow \{0, 1\}$, for all i , $c_i \leftarrow \text{Enc}_k(m_{b,i})$
3. $b' \leftarrow A(c_1, \dots, c_t)$

Adversary A **succeeds** if $b = b'$, and the experiment evaluates to **1** in this case.

Definition 3.4 Π is **multiple-message indistinguishable** if for **all PPT** attackers A , there is a **negligible** function ϵ such that

$$\Pr[\text{PrivK}_{A,\Pi}^{\text{mult}}(n) = 1] \leq 1/2 + \epsilon(n)$$

Q: Show that the pseudo OTP is **not** multiple-message indistinguishable



CPA-security

- We are **not** going to work with *multiple-message secrecy*



CPA-security

- We are **not** going to work with *multiple-message secrecy*

Instead, define something **stronger**: security against chosen-plaintext attacks (*CPA-security*)

- Nowadays, this is the **minimal** notion of security an encryption scheme should satisfy



CPA-security

- We are **not** going to work with *multiple-message secrecy*

Instead, define something **stronger**: security against chosen-plaintext attacks (*CPA-security*)

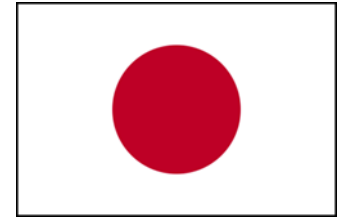
- Nowadays, this is the **minimal** notion of security an encryption scheme should satisfy

In practice, there are many ways an attacker can **influence** what gets encrypted

- Not clear how best to model
- Chosen-plaintext attacks encompasses any such influence



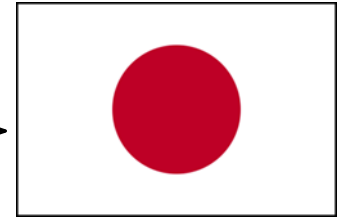
“Midway” story



“Midway” story



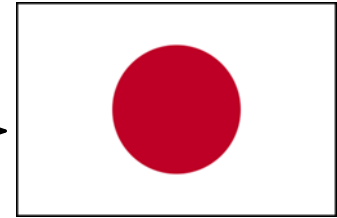
Will attack AF ..



“Midway” story



Will attack AF ..



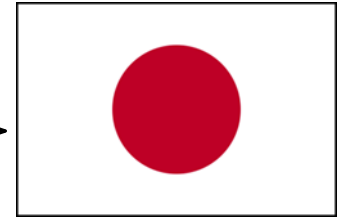
Help! Fresh water needed



“Midway” story



AF is short of water



Help! Fresh water needed



CPA-security

■ Fix Π , A

Define a randomized experiment $\text{PrivKCPA}_{A,\Pi}(n)$:

1. $k \leftarrow \text{Gen}(1^n)$
2. $A(1^n)$ **interacts** with an **encryption oracle** $\text{Enc}_k(\cdot)$, and then outputs m_0, m_1 of the same length
3. $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$, give c to A
4. A can **continue** to interact with $\text{Enc}_k(\cdot)$
5. A outputs b' ; A succeeds if $b = b'$, and experiment evaluates to 1 in this case



CPA-security

■ Fix Π , A

Define a randomized experiment $\text{PrivKCPA}_{A,\Pi}(n)$:

1. $k \leftarrow \text{Gen}(1^n)$
2. $A(1^n)$ **interacts** with an **encryption oracle** $\text{Enc}_k(\cdot)$, and then outputs m_0, m_1 of the same length
3. $b \leftarrow \{0, 1\}$, $c \leftarrow \text{Enc}_k(m_b)$, give c to A
4. A can **continue** to interact with $\text{Enc}_k(\cdot)$
5. A outputs b' ; A succeeds if $b = b'$, and experiment evaluates to 1 in this case

Definition 4.1 Π is **secure against chosen-plaintext attacks** (**CPA-secure**) if for **all PPT** attackers A , there is a **negligible** function ϵ such that

$$\Pr[\text{PrivKCPA}_{A,\Pi}(n) = 1] \leq 1/2 + \epsilon(n)$$



Impossible?

- Consider the following attacker A ;
 - Using a **chosen-plaintext attack**, get $c_0 = \text{Enc}_k(m_0)$ and $c_1 = \text{Enc}_k(m_1)$
 - Output m_0, m_1 ; get challenge ciphertext c
 - If $c = c_0$ output '0'; if $c = c_1$ output '1'
 - A succeeds with probability 1 (?)



Impossible?

- Consider the following attacker A ;
 - Using a **chosen-plaintext attack**, get $c_0 = \text{Enc}_k(m_0)$ and $c_1 = \text{Enc}_k(m_1)$
 - Output m_0, m_1 ; get challenge ciphertext c
 - If $c = c_0$ output '0'; if $c = c_1$ output '1'
 - A succeeds with probability 1 (?)
- This attack **only** works if encryption is deterministic!
 - **randomized** encryption must be used!
 - It really is a problem if an attacker can tell when **the same message is encrypted twice**

Pseudorandom functions

- Informally, a *pseudorandom function* “looks like” a random (i.e., uniform) function



Pseudorandom functions

- Informally, a *pseudorandom function* “looks like” a random (i.e., uniform) function

$Func_n =$ all functions mapping $\{0, 1\}^n$ to $\{0, 1\}^n$

How big is $Func_n$?

Pseudorandom functions

- Informally, a *pseudorandom function* “looks like” a random (i.e., uniform) function

$Func_n =$ all functions mapping $\{0, 1\}^n$ to $\{0, 1\}^n$

How big is $Func_n$?

- Can represent a function in $Func_n$ using $n \cdot 2^n$ bits



Pseudorandom functions

- Informally, a *pseudorandom function* “looks like” a random (i.e., uniform) function

$Func_n$ = all functions mapping $\{0, 1\}^n$ to $\{0, 1\}^n$

How big is $Func_n$?

- Can represent a function in $Func_n$ using $n \cdot 2^n$ bits

$$\Rightarrow |Func_n| = 2^{n \cdot 2^n}$$



Pseudorandom functions

- Informally, a *pseudorandom function* “looks like” a random (i.e., uniform) function

$Func_n$ = all functions mapping $\{0, 1\}^n$ to $\{0, 1\}^n$

How big is $Func_n$?

- Can represent a function in $Func_n$ using $n \cdot 2^n$ bits

$$\Rightarrow |Func_n| = 2^{n \cdot 2^n}$$

Q: how many functions are there mapping from $\{0, 1\}^n$ to $\{0, 1\}^m$?



Random functions vs. pseudorandom functions

- Choose uniform $f \in \text{Func}_n$



Random functions vs. pseudorandom functions

- Choose uniform $f \in \text{Func}_n$
- **Equivalent:** for each $x \in \{0, 1\}^n$, choose $f(x)$ **uniformly** in $\{0, 1\}^n$
 - I.e., fill up the function table with uniform values

Random functions vs. pseudorandom functions

- Choose uniform $f \in \text{Func}_n$
- **Equivalent:** for each $x \in \{0, 1\}^n$, choose $f(x)$ **uniformly** in $\{0, 1\}^n$
 - I.e., fill up the function table with uniform values
- Informally, a *pseudorandom function* “looks like” a random function
 - It does **not** make sense to talk about any **fixed** function being pseudorandom. We look instead at *keyed* functions



Keyed functions

- Let $F : \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$ be an efficient, deterministic algorithm
 - Define $F_k(x) = F(k, x)$
 - The first input is called the *key*



Keyed functions

- Let $F : \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$ be an efficient, deterministic algorithm
 - Define $F_k(x) = F(k, x)$
 - The first input is called the *key*
- Assume that F is *length preserving*: $F(k, x)$ only defined if $|k| = |x|$, in which case, $|F(k, x)| = |k| = |x|$



Keyed functions

- Let $F : \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$ be an efficient, deterministic algorithm
 - Define $F_k(x) = F(k, x)$
 - The first input is called the *key*
- Assume that F is *length preserving*: $F(k, x)$ only defined if $|k| = |x|$, in which case, $|F(k, x)| = |k| = |x|$
- Choosing a *uniform* $k \in \{0, 1\}^n$ is equivalent to choosing the function $F_k : \{0, 1\}^n \rightarrow \{0, 1\}^n$
 - I.e., for fixed key length n , the algorithm F defines a *distribution* over functions in $Func_n$!



Keyed functions

- Let $F : \{0, 1\}^* \times \{0, 1\}^* \rightarrow \{0, 1\}^*$ be an efficient, deterministic algorithm
 - Define $F_k(x) = F(k, x)$
 - The first input is called the *key*
- Assume that F is *length preserving*: $F(k, x)$ only defined if $|k| = |x|$, in which case, $|F(k, x)| = |k| = |x|$
- Choosing a *uniform* $k \in \{0, 1\}^n$ is equivalent to choosing the function $F_k : \{0, 1\}^n \rightarrow \{0, 1\}^n$
 - I.e., for fixed key length n , the algorithm F defines a *distribution* over functions in $Func_n$!
 - E.g., $F(k, x) = k$, $F(k, x) = k \oplus x$



- The number of functions in $Func_n$ is $2^{n \cdot 2^n}$



- The number of functions in $Func_n$ is $2^{n \cdot 2^n}$
- $\{F_k\}_{k \in \{0,1\}^n}$ is a subset of $Func_n$
 - The number of functions in $\{F_k\}_{k \in \{0,1\}^n}$ is **at most 2^n**

- The number of functions in $Func_n$ is $2^{n \cdot 2^n}$
- $\{F_k\}_{k \in \{0,1\}^n}$ is a subset of $Func_n$
 - The number of functions in $\{F_k\}_{k \in \{0,1\}^n}$ is **at most 2^n**

Definition 4.2 F is a *pseudorandom function* if F_k , for uniform $k \in \{0,1\}^n$ is **indistinguishable** from a uniform function $f \in Func_n$.
Formally, for **all** poly-time distinguishers D :

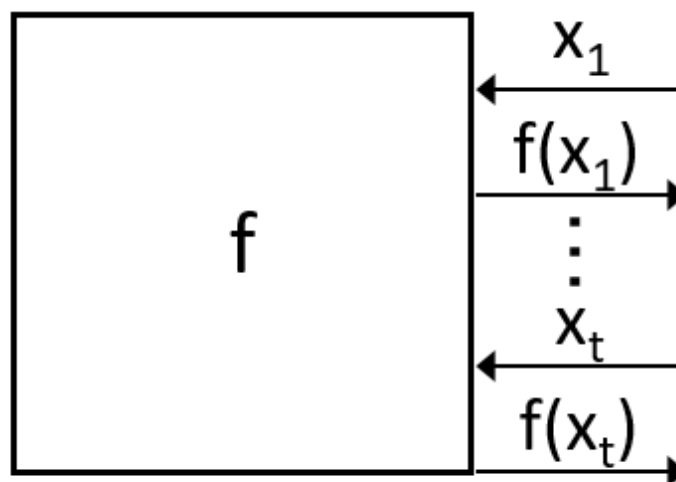
$$|\Pr_{k \leftarrow \{0,1\}^n}[D^{F_k(\cdot)}(1^n) = 1] - \Pr_{f \leftarrow Func_n}[D^{f(\cdot)}(1^n) = 1]| \leq \epsilon(n)$$



PRFs

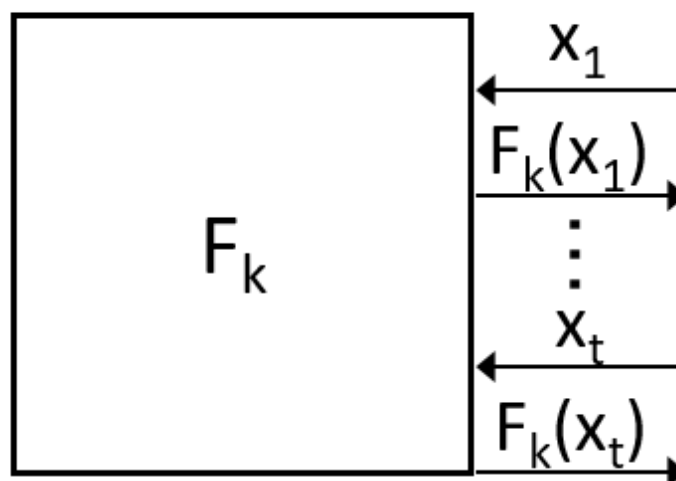
$f \in \text{Func}_n$ chosen
uniformly at random

World 0



World 1

$k \in \{0,1\}^n$ chosen
uniformly at random



??



(poly-time)

Pseudorandom permutations (PRPs)

- Let $f \in \text{Func}_n$



Pseudorandom permutations (PRPs)

- Let $f \in \text{Func}_n$
 f is a *permutation* if it is a bijection
 - This means that the *inverse* f^{-1} exists



Pseudorandom permutations (PRPs)

- Let $f \in \text{Func}_n$
 f is a *permutation* if it is a bijection
 - This means that the *inverse* f^{-1} exists
- Let $\text{Perm}_n \subset \text{Func}_n$ be the set of permutations
 - What is $|\text{Perm}_n|$?



Pseudorandom permutations

- Let F be a length-preserving, keyed function



Pseudorandom permutations

- Let F be a length-preserving, keyed function
- F is a *keyed permutation* if
 - F_k is a permutation for every k
 - F_k^{-1} is *efficiently computable* (where $F_k^{-1}(F_k(x)) = x$)



Pseudorandom permutations

- Let F be a length-preserving, keyed function
- F is a *keyed permutation* if
 - F_k is a permutation for every k
 - F_k^{-1} is *efficiently computable* (where $F_k^{-1}(F_k(x)) = x$)
- **Definition 4.3** F is a *pseudorandom permutation* if F_k , for **uniform** key $k \in \{0, 1\}^n$, is **indistinguishable** from a uniform permutation $f \in \text{Perm}_n$



Pseudorandom permutations

- Let F be a length-preserving, keyed function
- F is a *keyed permutation* if
 - F_k is a permutation for every k
 - F_k^{-1} is *efficiently computable* (where $F_k^{-1}(F_k(x)) = x$)
- **Definition 4.3** F is a *pseudorandom permutation* if F_k , for **uniform** key $k \in \{0, 1\}^n$, is **indistinguishable** from a uniform permutation $f \in \text{Perm}_n$
- For large enough n , a random permutation is **indistinguishable** from a random function.
 - In practice, PRPs are also good PRFs



PRFs vs. PRGs

- PRF F immediately implies a PRG G :
 - Define $G(k) = F_k(0 \dots 0) | F_k(0 \dots 1)$
 - I.e., $G(k) = F_k(\langle 0 \rangle) | F_k(\langle 1 \rangle) | F_k(\langle 2 \rangle) | \dots$,
where $\langle i \rangle$ denotes the n -bit encoding of i



PRFs vs. PRGs

- PRF F immediately implies a PRG G :
 - Define $G(k) = F_k(0 \dots 0) | F_k(0 \dots 1)$
 - I.e., $G(k) = F_k(\langle 0 \rangle) | F_k(\langle 1 \rangle) | F_k(\langle 2 \rangle) | \dots$,
where $\langle i \rangle$ denotes the n -bit encoding of i
- PRF can be viewed as a PRG with random access to **exponentially** long output
 - The function F_k can be viewed as the $n2^n$ -bit string $F_k(0 \dots 0) | \dots | F_k(1 \dots 1)$

Do PRFs/PRPs exist?

- They are a stronger primitive than PRGs
 - though can be built from PRGs



Do PRFs/PRPs exist?

- They are a stronger primitive than PRGs
 - though can be built from PRGs

Theorem (Goldreich, Goldwasser, Micali 1984)

If the PRG Axiom is **true**, then there exist PRFs.

How to Construct Random Functions

ODED GOLDREICH, SHAFI GOLDWASSER,
AND SILVIO MICALI

Massachusetts Institute of Technology, Cambridge, Massachusetts

Abstract. A constructive theory of randomness for functions, based on computational complexity, is developed, and a pseudorandom function generator is presented. This generator is a deterministic polynomial-time algorithm that transforms pairs (g, r) , where g is *any* one-way function and r is a random k -bit string, to polynomial-time computable functions $f_r: \{1, \dots, 2^k\} \rightarrow \{1, \dots, 2^k\}$. These f_r 's cannot be distinguished from *random* functions by any probabilistic polynomial-time algorithm that asks and receives the value of a function at arguments of its choice. The result has applications in cryptography, random constructions, and complexity theory.

Categories and Subject Descriptors: F.0 [Theory of Computation]: General; F.1.1 [Computation by Abstract Devices]: Models of Computation—*computability theory*; G.0 [Mathematics of Computing]: General; G.3 [Mathematics of Computing]: Probability and Statistics—*probabilistic algorithms; random number generation*

General Terms: Algorithms, Security, Theory

Additional Key Words and Phrases: Cryptography, one-way functions, prediction problems, randomness

I have set up on a Manchester computer a small programme using only 1000 units of storage, whereby the machine supplied with one sixteen figure number replies with another within two seconds. I would defy anyone to learn from these replies sufficient about the programme to be able to predict any replies to untried values.

A. TURING

Do PRFs/PRPs exist?

- They are a stronger primitive than PRGs
 - though can be built from PRGs
- In practice, **block ciphers** are used



Next Lecture

- block cipher ...

